

Search— Heuristic Search

KUO-SHIH TSENG (曾國師)
DEPARTMENT OF MATHEMATICS
NATIONAL CENTRAL UNIVERSITY, TAIWAN

2020/03/10

Course Announcement

- If you enrolled in this course, you can access to Robotics Lab (M213) via your NCU ID since this week.
- Robotics Lab has
 - Minibot X 10
 - Turtlebot3 X10
 - Bebop X8
 - PC X1
 - Notebook X2
 - Monitor X4



Course Announcement

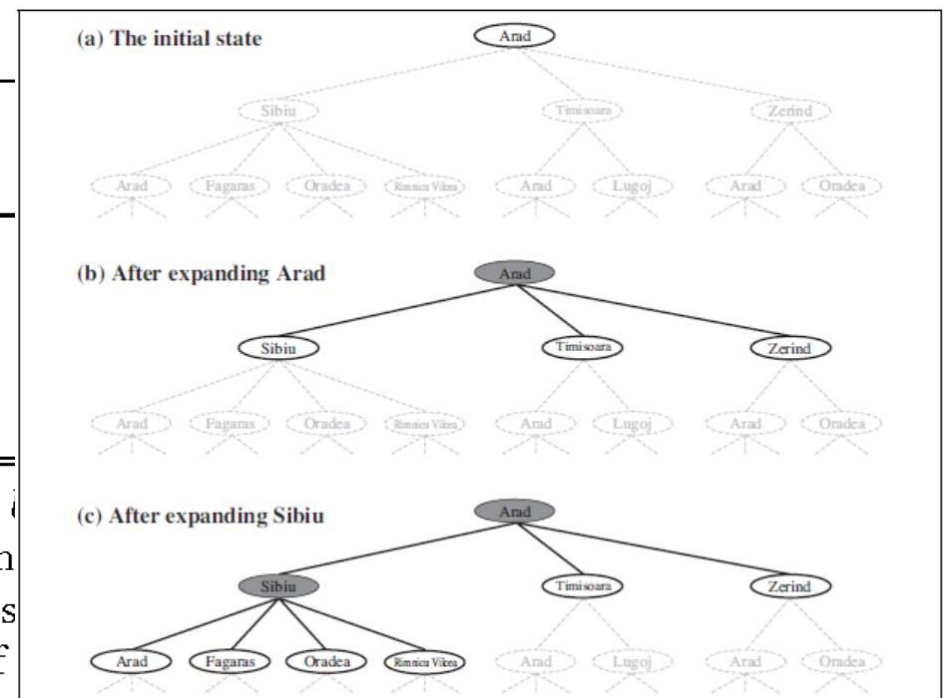
- HW1 was released on 3/02. Download it from eeclass.
- 1 programming problem and 2 theory problems.
- Deadline: 3/24(Wed.) 00:00am
- Delivery: Please update your HW to eecalss using electrical format. Compress your HW into a zip file including PDF and code.
- **Late policy:** If your HW is late for 1 day, the discount rate is 0.8. For 2 days, the discount rate is 0.8^2 . and so on.
- Start to work on it this week. You have 3 weeks to work.

Course Announcement

- DFS is not complete and optimal?

Criterion	Breadth-First	Uniform-Cost	Depth-First
Complete?	Yes ^a	Yes ^{a,b}	No
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$
Optimal?	Yes ^c	Yes	No

Figure 3.21 Evaluation of tree-search strategies. d is the depth of the shallowest solution; m is the maximum depth. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if ϵ is positive; ^c optimal if step costs are all identical; ^d if



Outline


- Informed search
 - Greedy best-first search
 - A* search
- Heuristic functions
- Online search
- Learning real-time A* (LRTA*)

Outline

[Problem solving]

Search problems


Adversarial Search



[Perception and Uncertainty]

Bayes Theorem

Bayes Filter and Smoothing






[Learning and Decision-making]

Supervised learning

Unsupervised learning

Reinforcement learning



Uninformed search and informed search

- Uninformed search:
 - There is **no** additional information beyond the given problem definition. It's also called blind search. The agent only can check if the current state is the goal state after actions.
- Informed search:
 - There is an additional information beyond the given problem definition. The agent can use this information to search more efficiently than blind search.

$$f(i) = \underbrace{g(i)}_{past} + \underbrace{h(i)}_{future}$$

f : evaluation function

g : cost function

h : heuristic function

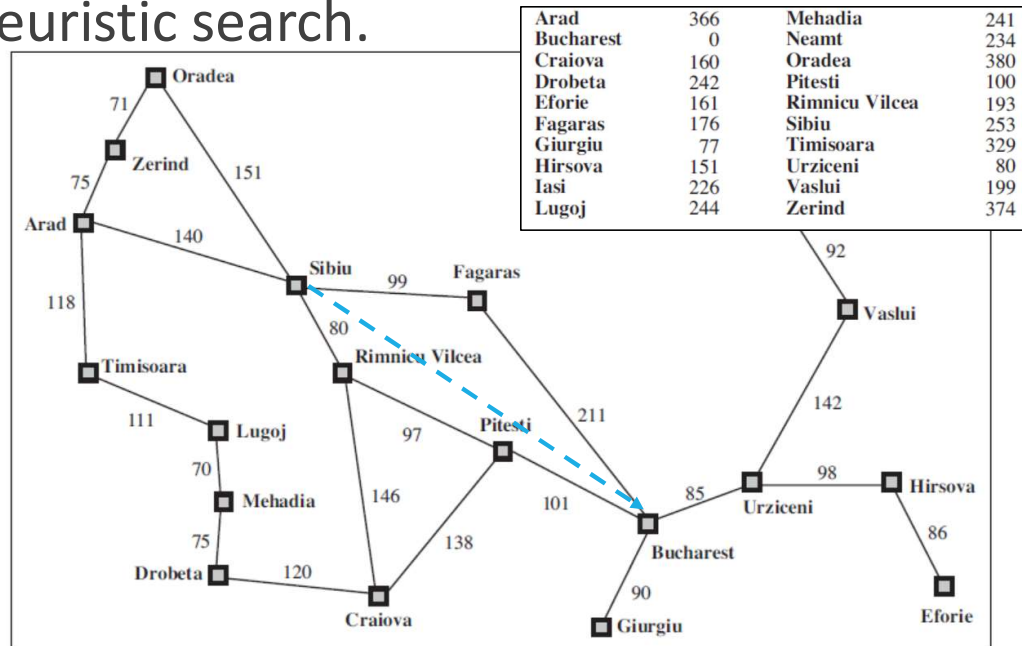
$$\begin{cases} \text{Uninformed search : } f(i) = g(i) \\ \text{Heuristic search : } f(i) = g(i) + h(i) \end{cases}$$

Informed search

- There is an additional information beyond the given problem definition. For example, in Romania-distance problem, the estimated cost to the goal is $h(i)$, where i is the current node. Informed search is also called heuristic search.

- h : Heuristic function
- Algorithms:
 - Greedy best-first search
 - A* search


* Heuristic in Greek means “to find.”




Informed search— Greedy best-first search

- Greedy best-first search: Expand the node with the lowest path cost from the current node to goal (heuristic function).

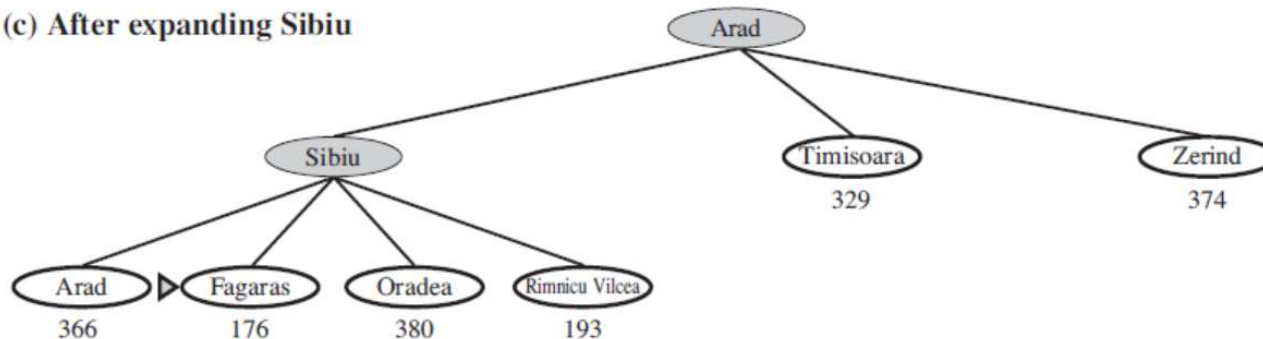
(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu

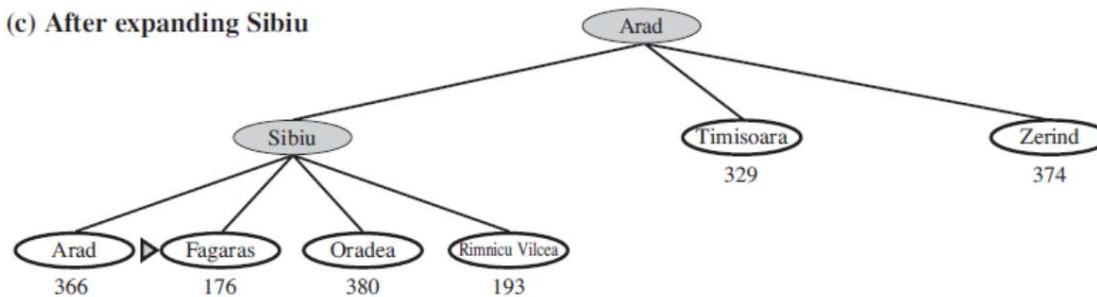


Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Informed search— Greedy best-first search

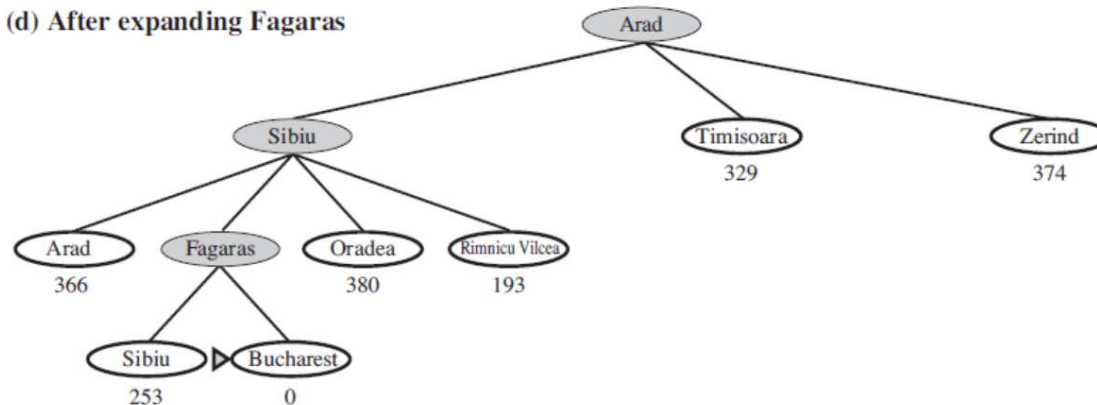
- Greedy best-first search: Expand the node with the lowest path cost from the current node to goal (heuristic function).

(c) After expanding Sibiu



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(d) After expanding Fagaras



Informed search— Greedy best-first search

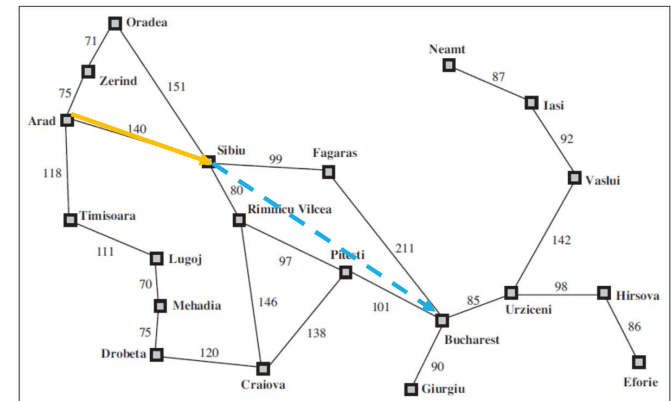
- Greedy best-first search is similar to uniform cost search but it use h function instead of g function. The data structure of greedy best-first search is **priority queue**.
- This approach is not optimal and not complete.
- UCS considers the **path cost** (g) while Greedy best-first search considers the **future cost** (h, heuristic). Could we consider both of them to find a better solution?
- How about Heuristic search : $f(i) = g(i) + h(i)$?

Informed search— A*

- A*: Expand the node with the lowest evaluation cost.
- $f(i) = g(i) + h(i)$
- $f(i)$: Evaluation cost
- $g(i)$: distance from start to current node i
- $h(i)$: *estimated* distance from goal to current node i

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374


$$f(i) = \underbrace{g(i)}_{past} + \underbrace{h(i)}_{future}$$




Informed search— A*

- A*: Expand the node with the lowest evaluation cost.

(a) The initial state

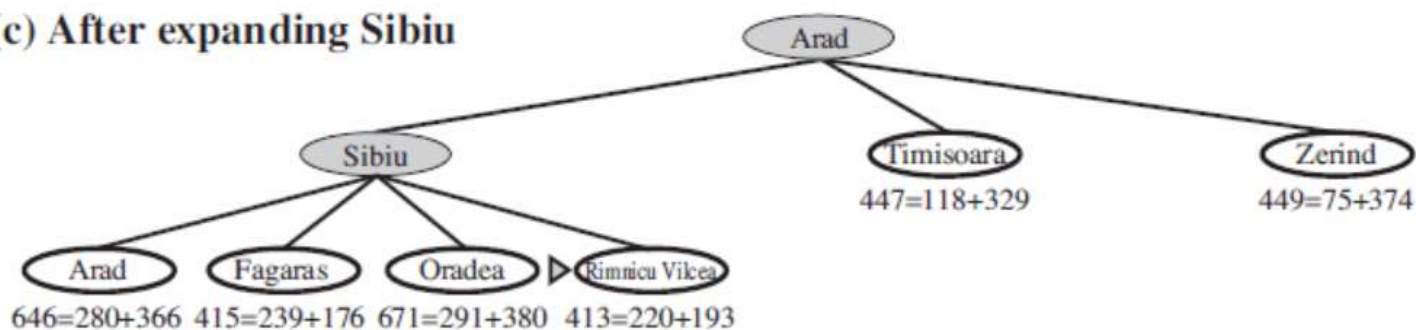


(b) After expanding Arad



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

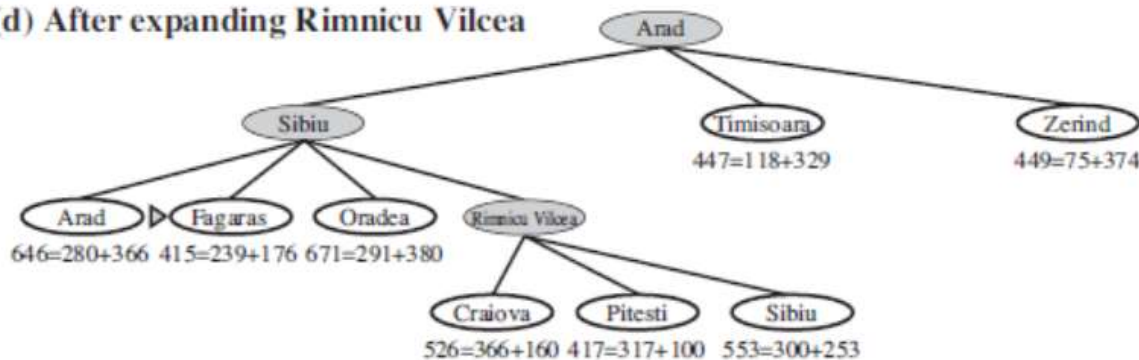
(c) After expanding Sibiu



Informed search— A*

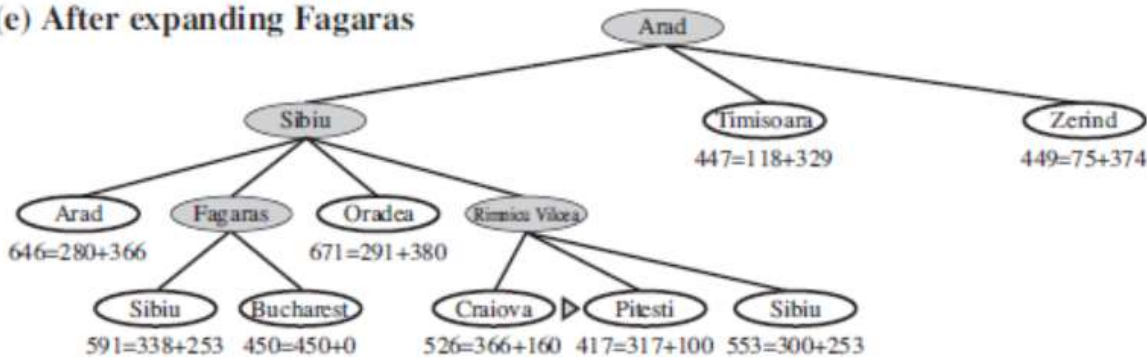
- A*: Expand the node with the lowest evaluation cost.

(d) After expanding Rimnicu Vilcea



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

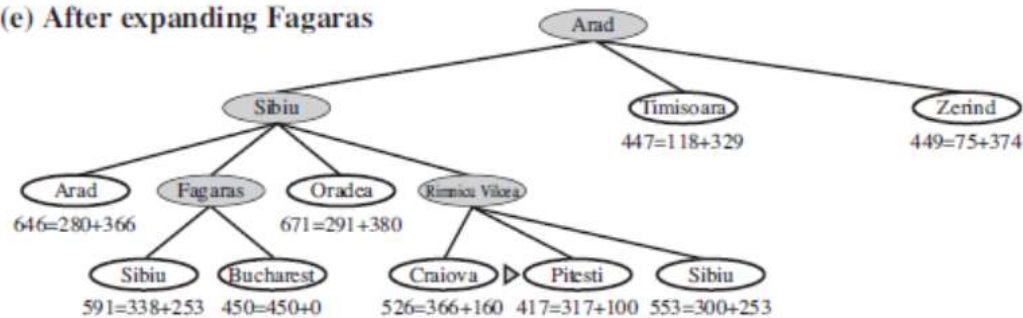
(e) After expanding Fagaras



Informed search— A*

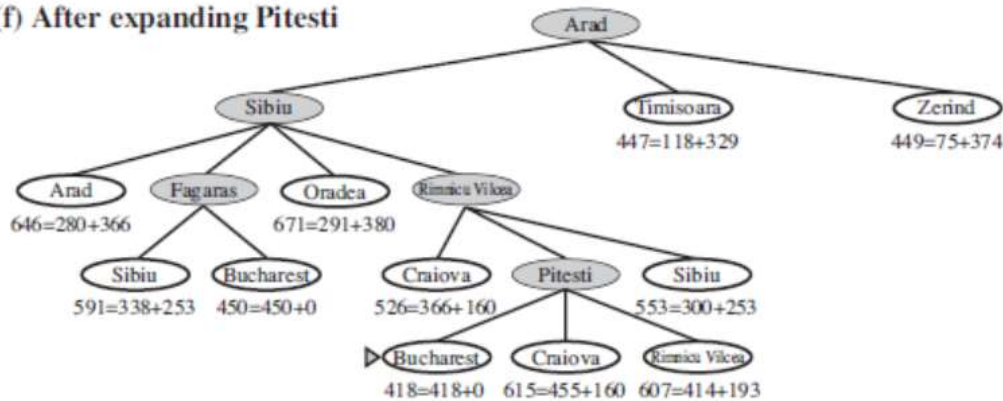
- A*: Expand the node with the lowest evaluation cost.

(e) After expanding Fagaras



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(f) After expanding Pitesti



Informed search— A*

- Conditions for optimality: Admissibility and Consistency
- Admissibility: An admissible heuristic is the one that ***never overestimates*** the cost to reach the goal.
- Consistency (monotonicity):

$$h(n) \leq c(n, a, n') + h(n')$$

where n' is the next state of n after action a

c is the cost function

The proof of A* optimality can be found in AIMA P.95.

Heuristic functions

- How to find a heuristic function is the key of heuristic search.
- Different heuristic functions could generate different search tree. We hope the number of branches is fewer.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

h_1 = # of misplaced tiles

h_2 = sum of the distance of the tiles from their goal positions.

(Manhattan distance)

$$h_1 = 8$$

$$h_2 = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$$

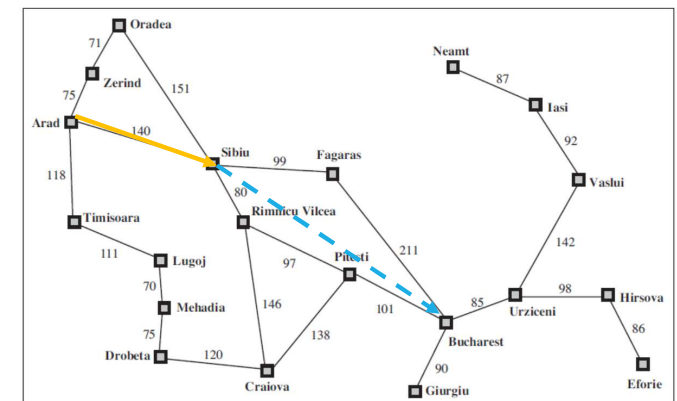
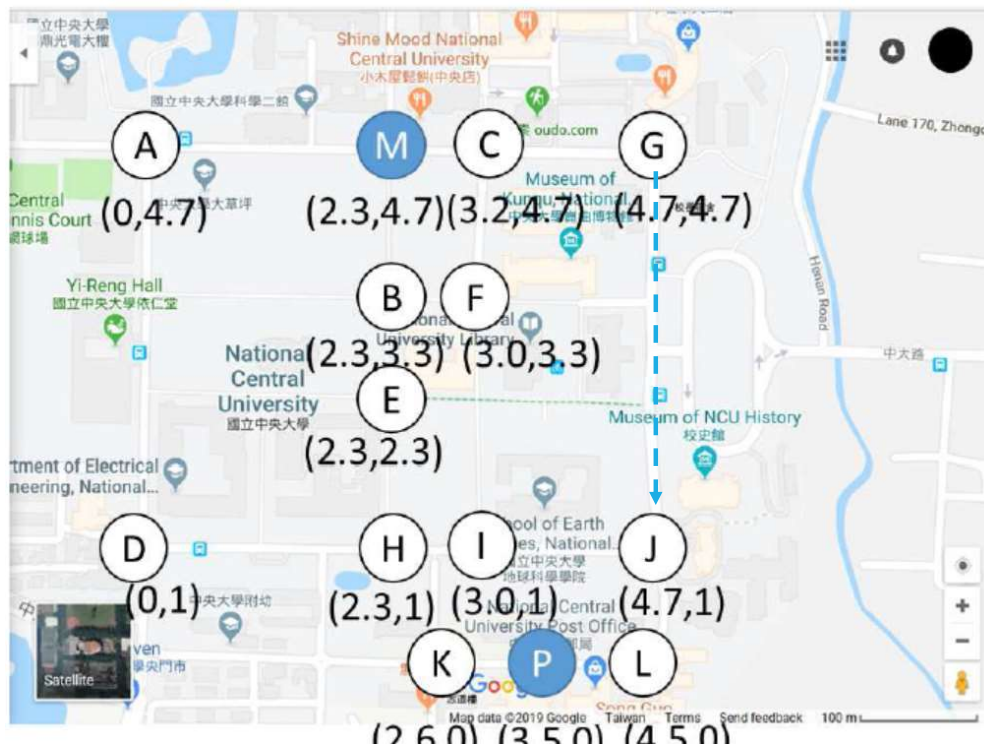
Heuristic functions

d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Figure 3.29 Comparison of the search costs and effective branching factors for the ITERATIVE-DEEPENING-SEARCH and A^* algorithms with h_1 , h_2 . Data are averaged over 100 instances of the 8-puzzle for each of various solution lengths d .

Heuristic functions

- Admissibility: An admissible heuristic is the one that ***never overestimates*** the cost to reach the goal.



Online Search

- The aforementioned search is offline search, which the agent has the environment information and searches for solutions.
- Online search is the case for unknown environments, where the robot doesn't know
 1. What the state is. $S : state \rightarrow ???$
 2. What the outcome of an action. $s, a \rightarrow ???$
- The robot needs to explore the environments and learn to find optimal actions via interaction.
- Online search is a deterministic case of *reinforcement learning*.

Online Search

- Online search is that the robot first takes an action, then it observes the environment and compute the next action.
- The robot does trial-and-error actions to improve its solutions.

[*GIVEN*]


A: action

$c(s, a, s')$: step – cost function

Goal – Test(*s*)

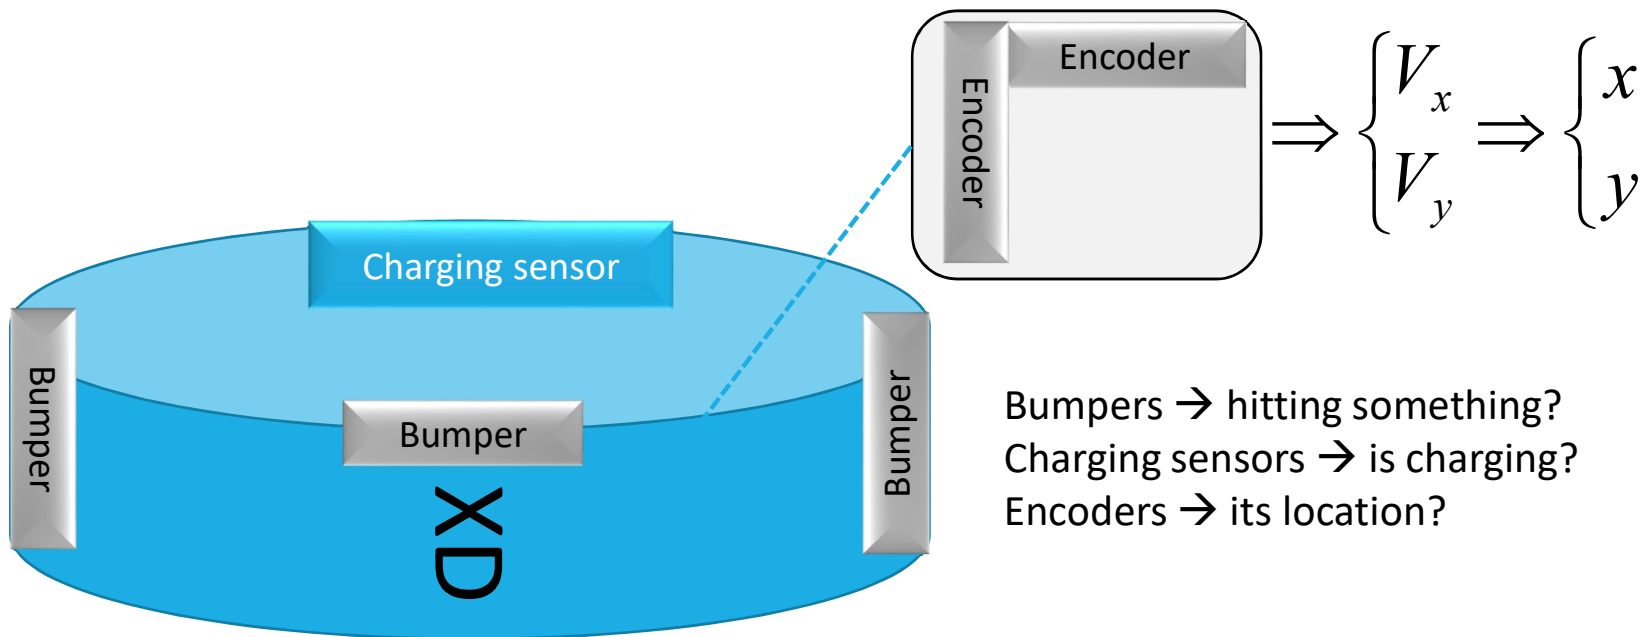
[Find]

The optimal (lowest cost) path




Online Search

- There is a cleaner robot with bumpers, charging sensors and encoders. Its goal is to move to the charging station with the lowest cost.



Online Search

- The robot does trial-and-error actions to improve its solutions.
- Why does the robot need to try?



3				Goal
2				+10
1	START			
	1	2	3	4

3				Goal
2				+10
1	START			
	1	2	3	4

3				Goal
2				+10
1	START			
	1	2	3	4

It needs to learn to work in different environments!

Online Search

- There is a cleaner robot with bumpers, charging sensors and encoders. Its goal is to move to the charging station with the lowest cost.

[GIVEN]

$A: action = \{left, up, right\}$


$c(s, a, s') : \text{step - cost function (Encoders, bumpers)}$

$\begin{cases} \text{hit wall} \rightarrow +2 \\ \text{hit } (4,2) \rightarrow +10 \\ \text{others} \rightarrow +1 \end{cases}$

$Goal - Test(s)$

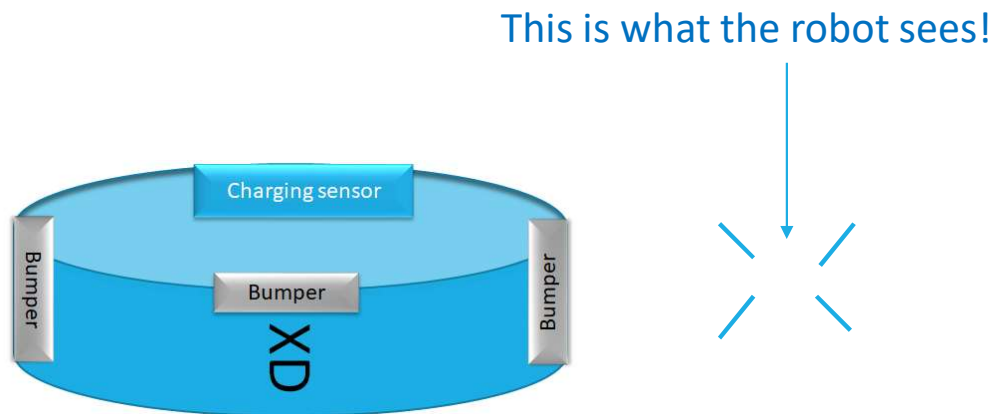
(charging sensors)

[Environment]



Online Search

- There is a cleaner robot with bumpers, charging sensors and encoders. Its goal is to move to the charging station with the lowest cost.




This is what you see!

3			Goal	
2			+10	
1	START			
	1	2	3	4


Online Search

- Illustration of the robot's world




Online Search

- The robot has to learn the cost via sequential actions.



$$c(s, a, s') = 2 \quad c(s, a, s') = 1 \quad c(s, a, s') = 1$$

(This is just an illustration of offline search. In online search, the robot cannot know the children nodes without experience.)



We want to apply A* to online search. The robot needs to **remember** its experience.


$$f(i) = \underbrace{g(i)}_{past} + \underbrace{h(i)}_{future}$$

If it's a offline search. We can build a tree or graph and then apply A* to find the optimal actions.

h : manhattan distance

LRTA*

- Let's think about how to revise A* for online search
- For offline search, the robot has the environment information. However, for online search, the robot doesn't know how large the environment is. Hence, the robot needs to adopt a memory efficiency way – Markov chain.
- The state (s') at time $t+1$ only depends on the state (s) at time t .

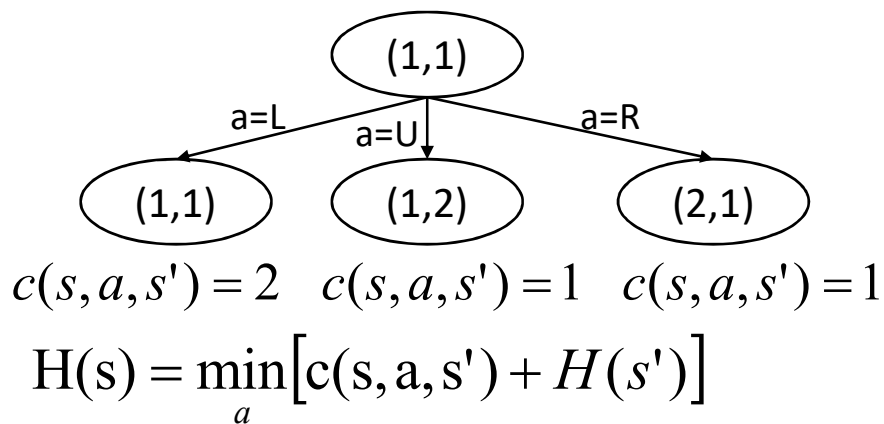


$$f(i) = \underbrace{g(i)}_{past} + \underbrace{h(i)}_{future}$$

$$H(s) = \underbrace{c(s, a, s')}_{past} + \underbrace{h(s')}_{future}$$


LRTA*

- Assume the robot has a heuristic function. The robot should choose the action with the lowest value of evaluation function



$$H(s) = \min_a \begin{cases} 2 + 5 = 7 & \text{(Left)} \\ 1 + 4 = 5 & \text{(Up)} \\ 1 + 4 = 5 & \text{(Right)} \end{cases}$$

The robot should choose Up or Right. However, the assumption is the **cost** and **action outcome** are known. In online search, the robot needs to explore them!




LRTA*

- How could the robot explore the environment or act optimally based on an explored H table?
 - This is a trade-off between exploration and exploitation.
 - **Exploration** is to explore unknown states while **exploitation** is to utilize known knowledge.
1. ϵ -greedy was proposed to solve this issue. The robot will action randomly with ϵ probability and select best action with $(1 - \epsilon)$ probability.
 2. Another way is to explore everywhere. Set unexplored $H(s)$ as a small value.

LRTA*

- The robot needs to **remember** its past experience to improve its actions.



$$[s = (1,1), a = L] \rightarrow s' = (1,1)$$

1. **remember** the result table.
2. **remember** the estimated cost table.

The result table tells the robot the next state (s') given (s, a) while the H table tells the robot what's the estimated cost at s .

Then, the robot can learn to find a solution!

It's called LRTA*[1], which is a special case of reinforcement learning.

LRTA* is the bridge between classic AI and modern AI.

[1] R.E. Korf, "Real-time heuristic search," Artificial Intelligence, 1990.

LRTA*

function LRTA*-AGENT(s') **returns** an action

inputs: s' , a percept that identifies the current state

persistent: *result*, a table, indexed by state and action, initially empty

H, a table of cost estimates indexed by state, initially empty

s, *a*, the previous state and action, initially null

if GOAL-TEST(s') **then return** *stop*

if s' is a new state (not in *H*) **then** $H[s'] \leftarrow h(s')$ (for exploration)

if *s* is not null

$result[s, a] \leftarrow s'$

$H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$

$a \leftarrow$ an action *b* in *ACTIONS*(s') that minimizes $LRTA^*-COST(s', b, result[s', b], H)$


$s \leftarrow s'$

return *a*

function LRTA*-COST(*s*, *a*, s' , *H*) **returns** a cost estimate

if s' is undefined **then return** $h(s)$ Set heuristic as estimated cost (for exploration)

else return $c(s, a, s') + H[s']$ (past cost + Heuristic)



[Past]


Update result table

Update H table

[Future]

Select action

LRTA*



$$c(s, a, s') = 2$$

$$H[s' = (1,1)] = 5$$

$$s' = (1,1) \leftarrow [s = (1,1), a = L]$$

$$\begin{cases} b = L \rightarrow H[s = (1,1)] = 2 + 5 \\ b = U \rightarrow \text{undefined}, h[s = (1,1)] = 5 \\ b = R \rightarrow \text{undefined}, h[s = (1,1)] = 5 \end{cases}$$

$$H[s = (1,1)] \leftarrow 5$$

$$a \leftarrow U$$


$$s \leftarrow s' = (1,1)$$



if s is not null
 $result[s, a] \leftarrow s'$
 $H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$
 $a \leftarrow$ an action b in $ACTIONS(s')$ that minimizes $LRTA^*-COST(s', b, result[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow H[s'' = (1,1)] = 7 \\ b = U \rightarrow \text{undefined}, h[s' = (1,1)] = 5 \\ b = R \rightarrow \text{undefined}, h[s' = (1,1)] = 5 \end{cases}$$

LRTA*



$$c(s, a, s') = 1$$

$$H[s' = (1,2)] = 4$$


$$s' = (1,2) \leftarrow [s = (1,1), a = U]$$

$$\begin{cases} b = L \rightarrow H[s = (1,1)] = 2 + 5 \\ b = U \rightarrow H[s = (1,2)] = 1 + 4 = 5 \\ b = R \rightarrow \text{undefined}, h[s = (1,1)] = 5 \end{cases}$$

$$H[s = (1,1)] \leftarrow 5$$

$$a \leftarrow L$$

$$s \leftarrow s' = (1,2)$$



if s is not null


$$result[s, a] \leftarrow s'$$

$$H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$$

$a \leftarrow$ an action b in $ACTIONS(s')$ that minimizes $LRTA^*-COST(s', b, result[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow \text{undefined}, h[s' = (1,2)] = 4 \\ b = U \rightarrow \text{undefined}, h[s' = (1,2)] = 4 \\ b = R \rightarrow \text{undefined}, h[s' = (1,2)] = 4 \end{cases}$$

LRTA*



$$c(s, a, s') = 2$$


$$s' = (1,2) \leftarrow [s = (1,2), a = L]$$

$$\begin{cases} b = L \rightarrow H[s = (1,2)] = 2 + 4 \\ b = U \rightarrow \text{undefined}, h[s = (1,2)] = 4 \\ b = R \rightarrow \text{undefined}, h[s = (1,2)] = 4 \end{cases}$$

$$H[s = (1,2)] \leftarrow 4$$

$$a \leftarrow U$$


$$s \leftarrow s' = (1,2)$$



if s is not null
 $result[s, a] \leftarrow s'$
 $H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$
 $a \leftarrow$ an action b in $ACTIONS(s')$ that minimizes $LRTA^*-COST(s', b, result[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow 2 + h[s' = (1,2)] = 6 \\ b = U \rightarrow \text{undefined}, h[s' = (1,2)] = 4 \\ b = R \rightarrow \text{undefined}, h[s' = (1,2)] = 4 \end{cases}$$

LRTA*



$$c(s, a, s') = 1$$

$$H[s' = (1,3)] = h[s' = (1,3)] = 3$$


$$s' = (1,3) \leftarrow [s = (1,2), a = U]$$

$$\begin{cases} b = L \rightarrow H[s = (1,2)] = 2 + 4 \\ b = U \rightarrow H[s = (1,2)] = 1 + 4 \\ b = R \rightarrow \text{undefined}, h[s = (1,2)] = 4 \end{cases}$$

$$H[s = (1,2)] \leftarrow 4$$

$$a \leftarrow L$$

$$s \leftarrow s' = (1,3)$$



if s is not null


$$\text{result}[s, a] \leftarrow s'$$

$$H[s] \leftarrow \min_{b \in \text{ACTIONS}(s)} \text{LRTA}^*\text{-COST}(s, b, \text{result}[s, b], H)$$

$a \leftarrow$ an action b in $\text{ACTIONS}(s')$ that minimizes $\text{LRTA}^*\text{-COST}(s', b, \text{result}[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \\ b = U \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \\ b = R \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \end{cases}$$

LRTA*




$$s' = (1,3) \leftarrow [s = (1,3), a = L]$$

$$\begin{cases} b = L \rightarrow H[s = (1,3)] = 2 + 3 \\ b = U \rightarrow \text{undefined}, h[s = (1,3)] = 3 \\ b = R \rightarrow \text{undefined}, h[s = (1,3)] = 3 \end{cases}$$

$$H[s = (1,3)] \leftarrow 3$$

$$a \leftarrow U$$

$$s \leftarrow s' = (1,3)$$




```

if s is not null
    result[s, a] ← s'
    H[s] ← minb ∈ ACTIONS(s) LRTA*-COST(s, b, result[s, b], H)
    a ← an action b in ACTIONS(s') that minimizes LRTA*-COST(s', b, result[s', b], H)
    s ← s'
    
```

$$\begin{cases} b = L \rightarrow 2 + h[s' = (1,3)] = 5 \\ b = U \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \\ b = R \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \end{cases}$$

LRTA*




$$s' = (1,3) \leftarrow [s = (1,3), a = U]$$

$$\begin{cases} b = L \rightarrow H[s = (1,3)] = 2 + 3 \\ b = U \rightarrow H[s = (1,3)] = 2 + 3 \\ b = R \rightarrow \text{undefined}, h[s = (1,3)] = 3 \end{cases}$$

$$H[s = (1,3)] \leftarrow 3$$

$$a \leftarrow R$$


$$s \leftarrow s' = (1,3)$$



if s is not null
 $result[s, a] \leftarrow s'$
 $H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$
 $a \leftarrow$ an action b in $ACTIONS(s')$ that minimizes $LRTA^*-COST(s', b, result[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow 2 + h[s' = (1,3)] = 5 \\ b = U \rightarrow 2 + h[s' = (1,3)] = 5 \\ b = R \rightarrow \text{undefined}, h[s' = (1,3)] = 3 \end{cases}$$

LRTA*



$$H[s' = (2,3)] = h[s' = (2,3)] = 2$$


$$s' = (2,3) \leftarrow [s = (1,3), a = R]$$

$$\begin{cases} b = L \rightarrow H[s = (1,3)] = 2 + 3 \\ b = U \rightarrow H[s = (1,3)] = 2 + 3 \\ b = R \rightarrow H[s = (1,3)] = 1 + 2 \end{cases}$$

$$H[s = (1,3)] \leftarrow 3$$

$$a \leftarrow L$$

$$s \leftarrow s' = (2,3)$$



if s is not null


$$result[s, a] \leftarrow s'$$

$$H[s] \leftarrow \min_{b \in ACTIONS(s)} LRTA^*-COST(s, b, result[s, b], H)$$


$a \leftarrow$ an action b in $ACTIONS(s')$ that minimizes $LRTA^*-COST(s', b, result[s', b], H)$
 $s \leftarrow s'$

$$\begin{cases} b = L \rightarrow undefined, h[s' = (2,3)] = 2 \\ b = U \rightarrow undefined, h[s' = (2,3)] = 2 \\ b = R \rightarrow undefined, h[s' = (2,3)] = 2 \end{cases}$$


LRTA*



LRTA*



....




This is episode #1.

The robot built the H table and results table.

The robot can try more episodes to improve its actions.

Or use more exploration actions to find more possible solutions.



Properties of LRTA*

[Markov Chain]

In LRTA*, the s' only depends on s and a .

[Past + Future cost]

$$H(s) = c(s, a, s') + h(s')$$

A special case of Bellman equation

[Optimality]

Dynamic programming?

LRTA*

- Recalled optimality of A*.
- Conditions for optimality: Admissibility and Consistency
- Admissibility: An admissible heuristic is the one that ***never overestimates*** the cost to reach the goal.
- Consistency (monotonicity):

$$h(n) \leq c(n, a, n') + h(n')$$


where n' is the next state of n after action a

c is the cost function

LRTA* and Reinforcement Learning

- LRTA*

Deterministic action



$$s, a \rightarrow s'$$

L2: Uninformed search

L3: Heuristic search (LRTA*)


L4: Adversarial search

L5: Bayes theorem

L6: Bayes theorem over time

MDP (RL)

Probabilistic actions



$$P(s'|s, a)$$

L7: MDP

L9: Reinforcement learning

L10: GP and LWPR


L11: Naïve Bayes and Perceptron

L12: Adaboost

L13: Deep learning and DRL

POMDP

Probabilistic actions and states



$$P(s'|s, a), P(s)$$

L8: POMDP

Q&A



