

Markov Decision Process (MDP)

KUO-SHIH TSENG (曾國師)
DEPARTMENT OF MATHEMATICS
NATIONAL CENTRAL UNIVERSITY, TAIWAN

2021/04/07

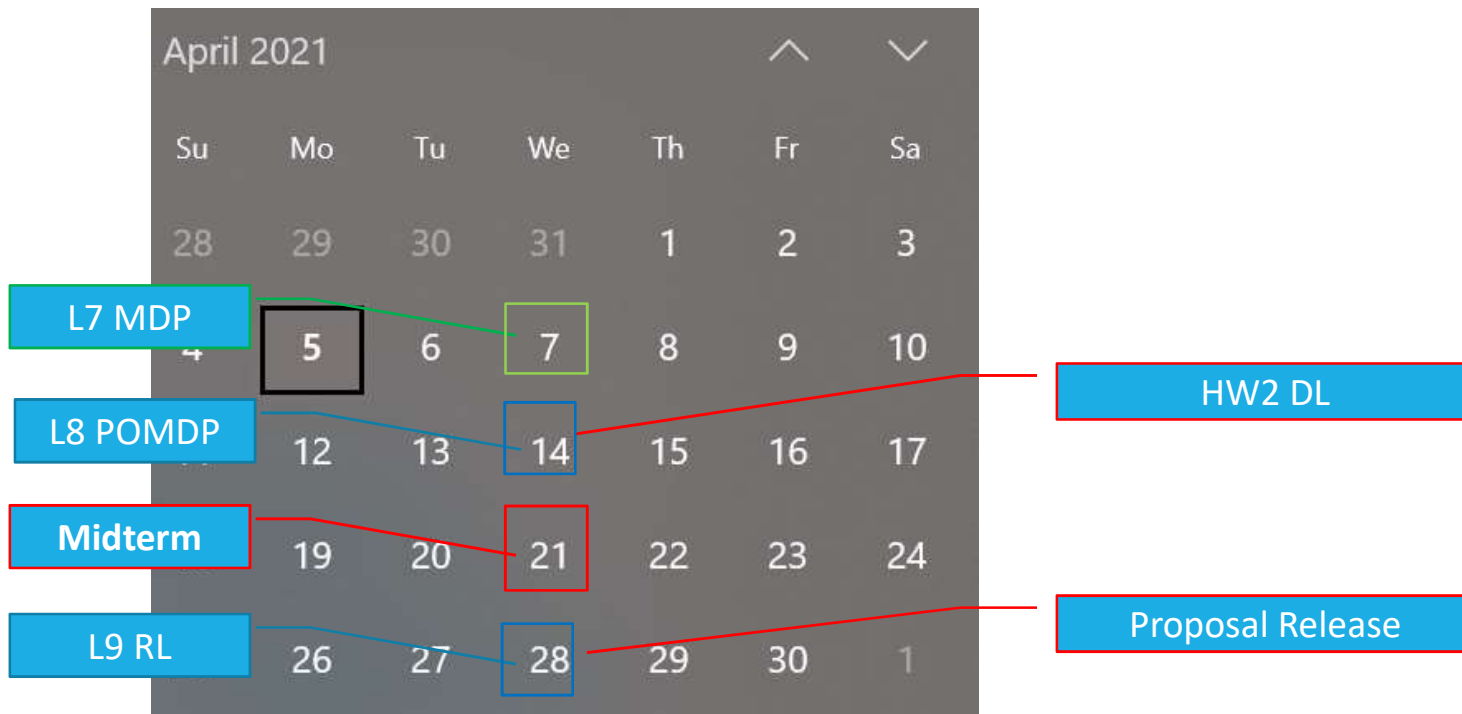
Course Announcement

- Work on your HW2 ASAP. The deadline is 4/14(Wed).
 - Bayesian inference (40%)
 - **MDP solver (40%)**
 - **A MDP problem (20%)**
- You can discuss course materials with me and classmates. You **should** work on your HW *independently*.
- NOTICE: I will **NOT** do curve fitting (e.g., “ $\sqrt{X} \cdot 6$ ” for your score) for your scores.
- **Late policy: If your HW is late for 1 day, the discount rate is 0.8. For 2 days, the discount rate is 0.8^2 . and so on.**
- Remember: life is a real-time model! Not a offline search!

Course Announcement

- ***Midterm (04/21/2020), 3-5pm, in M430***
 - Given a real world problem.
 - Design a perception and decision-making system for this problem using MDP , MCTS and Bayesian approaches.
- You can take one A4-size cheating sheet.
- You cannot use any electrical devices (e.g., Notebook or mobilephone), which can access to internet.
- You don't need calculators.
- You can find the **midterm_sample.pdf** on the eeclass.

Course Announcement



Course Announcement

May 2021

Su	Mo	Tu	We	Th	Fr	Sa
	3	4	5	6	7	8
	10	11	12	13	14	15
	17	18	19	20	21	22
	24	25	26	27	28	29
	31	1	2	3	4	5
	7	8	9	10	11	12
	14	15	16	17	18	19
	21	22	23	24	25	26

ROS
Tutorial

L11 NB
Perceptron

L12
Adaboost

L10 GP

L13
DL and DRL

L14
Kmeans, EM

Final Project
Presentation

Final Project
DEMO

HW3 released

Proposal DL

HW3 DL

Final Project report

Outline


- MDP
- Recitation: LRTA*
- Bellman equation
- Value iteration
- Policy iteration

Outline

[Problem solving]

Search problems


Adversarial Search



[Perception and Uncertainty]

Bayes Theorem

Bayes Filter and Smoothing






[Learning and Decision-making]

Supervised learning

Unsupervised learning

Reinforcement learning





MDP

- Why do humans buy lotteries?
- Why do humans keep doing something with low probability?
- Humans made decisions based on
 - Probability
 - Utility
 - **Future**

MDP

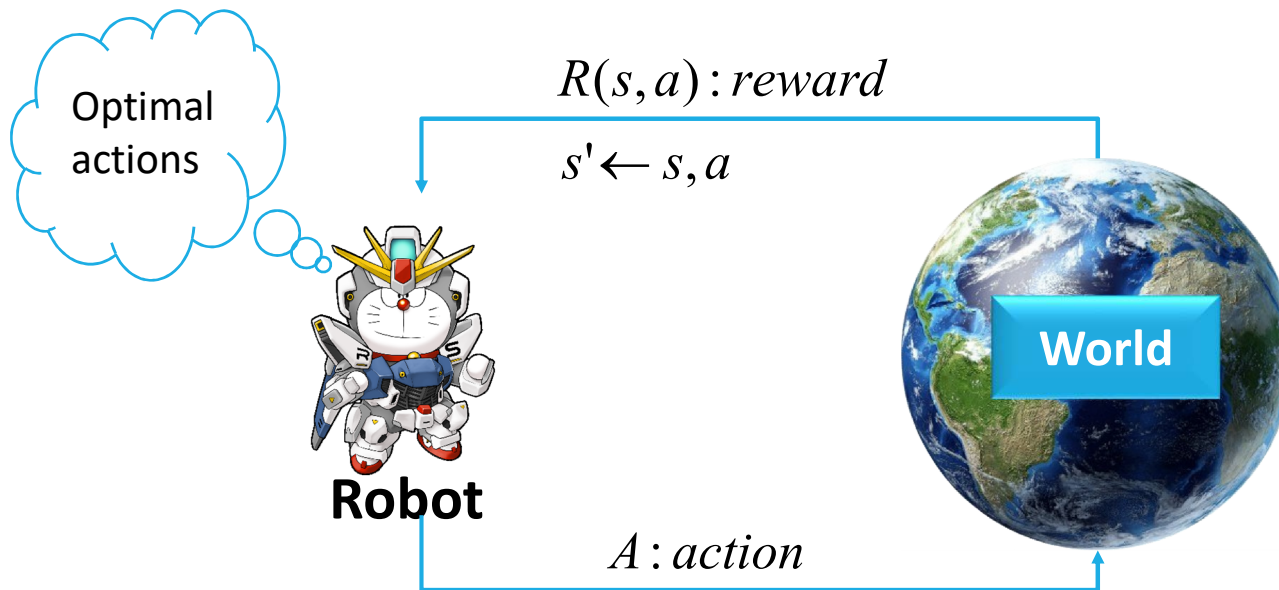
- Bayes decision only considers one-step expectation. For most of applications, the robot needs to do **sequential** optimal decisions for achieving the goal.
- In LRTA*, the transition/motion model is deterministic. Let's further consider a probabilistic transition/motion model for finding **sequential** optimal decisions.

$$\frac{p(open | z_2)}{p(close | z_2)} > \frac{(R_{CC} - R_{OC})}{(R_{OO} - R_{CO})} \Rightarrow \text{Decision : Move!}$$



MDP

- MDP is a model for finding sequential optimal decisions.
 - State: fully observable
 - State transition: stochastic **(Motion model)**



MDP

- MDP is a model for finding sequential optimal decisions.
 - State: fully observable
 - State transition: stochastic **(Motion model)**

[*GIVEN*]

S : state

A : action

$P(s'|s, a)$: Transition probability

$R(s, a)$: reward

γ : discount

s : state in t

s' : state in $t + 1$

[*Find*]

$\pi^* = \arg \max U^\pi(s)$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

π : policy

π^* : optimal policy

U : utility

MDP

- The M.S. life

[GIVEN]

$S : \text{state} \Rightarrow 4 \text{ states}$

$A : \text{action} \Rightarrow a \in \{\text{work hard}, \text{lazy}\}$

$P(s'|s, a) : \text{Transition probability}$

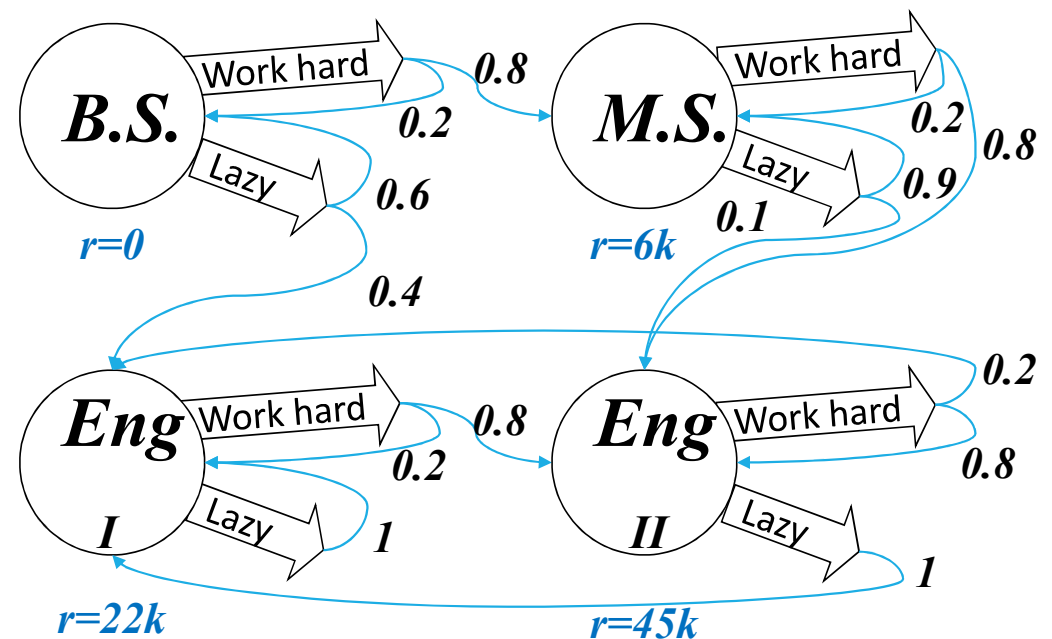
$R(s, a) : \text{reward}$

$\gamma : \text{discount} = 0.9$

[Find]

$\pi^* = \arg \max U^\pi(s)$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$



Stochastic automata (state machine) diagram

MDP

- 4X3 world

[GIVEN]

$S : state \Rightarrow (x, y), x \in \{1, \dots, 4\}, y \in \{1, \dots, 3\}$

$A : action \Rightarrow a \in \{\uparrow, \leftarrow, \rightarrow, \downarrow\}$

$P(s'|s, a) : Transition\ probability$


$R(s, a) : reward$

$\gamma : discount = 0.9$

[Find]

$\pi^* = \arg \max U^\pi(s)$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$




If the robot hits the wall,
It will bounce back.

$s=(4,3), R=+1$
 $s=(4,2), R=-1$
Others, $R=-0.04$

MDP


- Planning



How to solve these problems?

Recitation: LRTA*

- For offline search, the robot has the environment information. However, for online search, the robot doesn't know how large the environment is. Hence, the robot needs to adopt a memory efficiency way – Markov chain.
- The state (s') at time $t+1$ only depends on the state (s) at time t .



$$h[s' = (1,1)] = ?$$

$$f(i) = \underbrace{g(i)}_{past} + \underbrace{h(i)}_{future}$$

$$H(s) = \underbrace{c(s, a, s')}_{past} + \underbrace{h(s')}_{future}$$

(4,3)

Bellman Equation

- Assuming the agent chooses the optimal action,
The utility of a state is the immediate reward for that state + the expected discounted utility of the next state.
- Bellman equation is dynamic programming, which solving subproblem to find the optimal solution of the problem.

$$U(s) = \gamma \max_a \left[R(s, a) + \sum_{s'} U(s') P(s' | s, a) \right]$$

or

$$U(s) = \underbrace{R(s)}_{\text{immediate reward}} + \gamma \max_a \underbrace{\left[\sum_{s'} U(s') P(s' | s, a) \right]}_{\text{expected discounted utility}}$$

immediate reward

expected discounted utility

Bellman Equation

- To illustrate Bellman equation, let's look at an example in deterministic and probabilistic cases.
- A cleaner robot in a 3X3 world.
- State: $(1 \sim 3, 1 \sim 3)$
- Action: (left, up, right)
- Reward: +5 at goal (charging), -1 at other cells.
- Discount factor =1

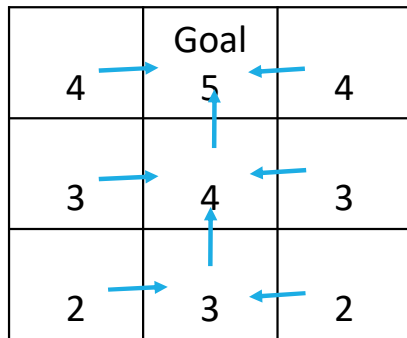
	Goal	

Bellman Equation

- Utility in deterministic cases

$$U(s) = R(s_0) + \gamma^1 R(s_1) + \gamma^2 R(s_2) + \dots + \gamma^n R(s_n)$$
$$= \sum_t \gamma^t R(s_t)$$

4	Goal 5	4
3	4	3
2	3	2



We can make optimal decisions with the utility function


4	Goal 5	4
3	4	3
2	3	2

We can compute the utility function from the goal to each state

Bellman Equation

- Utility in probabilistic cases
 - When the robot made an action, the next state s' is with uncertainty
- State: $(1 \sim 3, 1 \sim 3)$
- Action: (left, up, right)
- Transition probability: $P(s' | s, a)$

$$U(s) = E \left[\sum_{t=0}^n \gamma^t R(S_t) \right]$$




Bellman Equation

- Utility and action in deterministic cases
- After an action (a), $s \rightarrow s'$

$$U(s) = R(s) + \gamma \max U(s')$$

$$? = -1 + \max \begin{cases} U(s'_a) = 2 \\ U(s'_b) = 4 \\ U(s'_c) = 2 \end{cases}$$

4	Goal 5	4
3	s'_b 4	3
s'_{a2}		s'_{c2}

The robot chooses the next state with the maximal utility (optimal decision)


Bellman Equation

- Utility and action in probabilistic cases
- After an action (a), $s \rightarrow s'$ with $P(s' | s, a)$

$$U(s) = R(s) + \gamma \max_a \sum_{s'} U(s') P(s' | s, a)$$

$$U(s) = -1 + \max_a \begin{cases} 2 * 0.8 + 4 * 0.1 + 2 * 0.1 = 2.2 \\ 2 * 0.1 + 4 * 0.8 + 2 * 0.1 = 3.6 = 2.6 \\ 2 * 0.1 + 4 * 0.1 + 2 * 0.8 = 2.2 \end{cases}$$

Illustrate this case



Bellman equation shows the relationship between $U(s)$, $U(s')$ and actions (a).

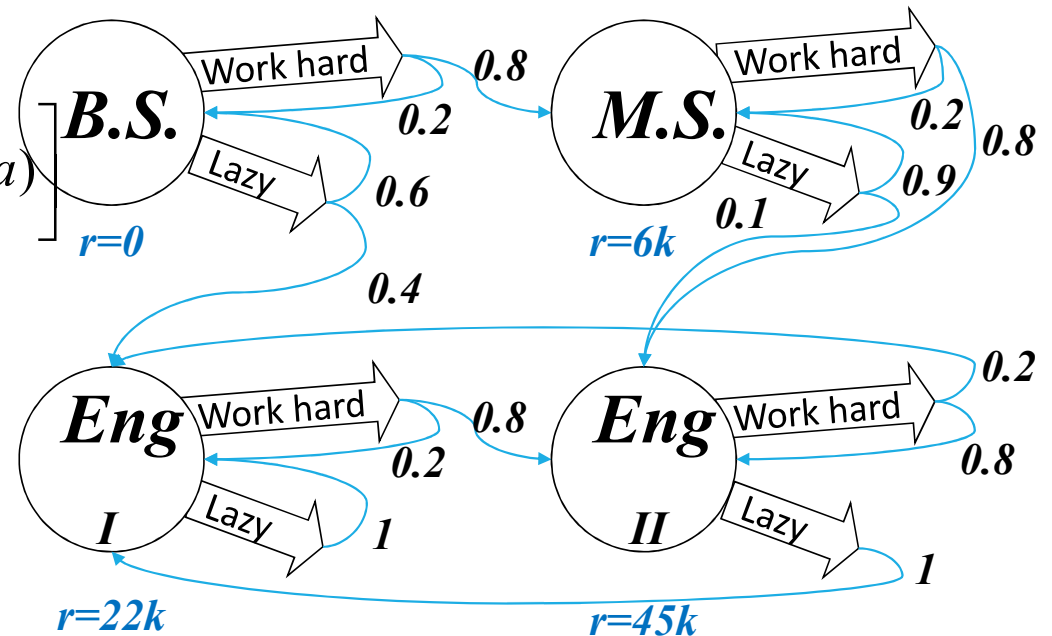
The assumption is “Markov chain.” The s' is only depending on s and a .

- Assuming the agent chooses the optimal action,
The utility of a state is the immediate reward for that state + the expected discounted utility of the next state.

Bellman Equation

- The M.S. life

$$U(s) = R(s) + \gamma \max_a \left[\sum_{s'} U(s') P(s' | s, a) \right]$$



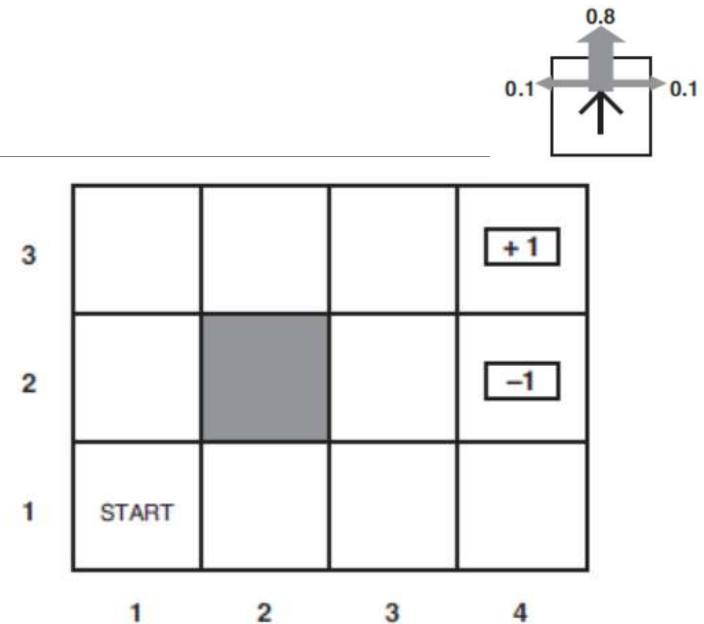
$$U(s = M.S.) = 6000 + \gamma \max_a \begin{bmatrix} 0.8U(s = E.II.) + 0.2U(s = M.S.) \\ 0.9U(s = M.S.) + 0.1U(s = E.II.) \end{bmatrix} \begin{matrix} \text{(work hard)} \\ \text{(lazy)} \end{matrix}$$

If we know U function, we can make optimal decisions! But, we don't know it.

Bellman Equation

- 4x3 world

$$U(s) = R(s) + \gamma \max_a \left[\sum_{s'} U(s') P(s' | s, a) \right]$$



$$U(1,1) = -0.04 + \gamma \max_a \begin{bmatrix} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1) \\ 0.9U(1,1) + 0.1U(1,2) \\ 0.9U(1,1) + 0.1U(2,1) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$

If we know U function, we can make optimal decisions! But, we don't know it.

Bellman Equation

- We try to find optimal actions for 1-step, 2-step to infinite-step.

- 1-step

$$\pi_1 = \arg \max U(s)$$

$$U_1(s) = \gamma R(s, a)$$

- 2-step

$$\pi_2 = \arg \max \left[R(s, a) + \sum_{s'} U_1(s') P(s' | s, a) \right]$$

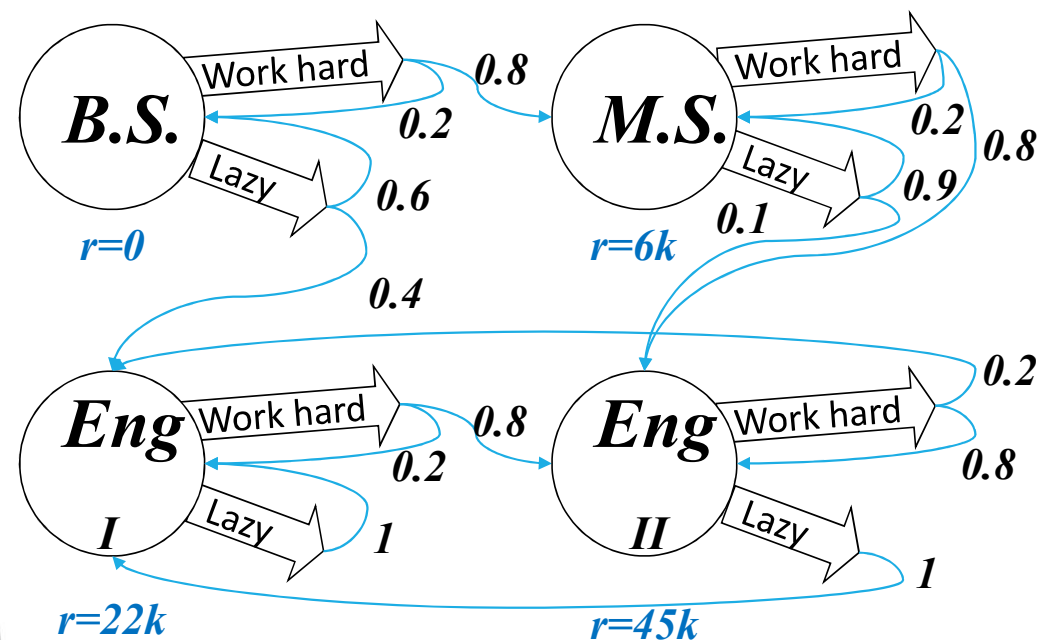
$$U_2(s) = \gamma \max \left[R(s, a) + \sum_{s'} U_1(s') P(s' | s, a) \right]$$

4	Goal 5	4
3	4	3
2	3	2

- T=1: greedy policy
- T>1: finite horizon case, typically no discount
- T=infty: infinite-horizon case, finite reward if discount < 1

Bellman Equation

- How about the utility in infinity steps?



Bellman Equation

- We try to find optimal actions for 1-step, 2-step to infinite-step.

- Nth-step

$$\pi_T = \arg \max \left[R(s, a) + \sum_{s'} U_{T-1}(s') P(s' | s, a) \right]$$

$$U_T(s) = \gamma \max \left[R(s, a) + \sum_{s'} U_{T-1}(s') P(s' | s, a) \right]$$

- Infinite-step

$$\pi_{\infty} = \arg \max \left[R(s, a) + \sum_{s'} U_{\infty}(s') P(s' | s, a) \right]$$


$$U_{\infty}(s) = \gamma \max \left[R(s, a) + \sum_{s'} U_{\infty}(s') P(s' | s, a) \right]$$

4	Goal 5	4
3	4	3
2	3	2

- T=1: greedy policy
- T>1: finite horizon case, typically no discount
- T=infty: infinite-horizon case, finite reward if discount < 1

Value iteration


- If there are N states, there are N Bellman equations. There are N unknown utilities of states.
- However, we cannot use $A=BX$ to solve this problem since it's not a linear algebra formulation. (max is a nonlinear operator)




$$U(1,1) = -0.04 + \gamma \max_a \begin{bmatrix} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1) \\ 0.9U(1,1) + 0.1U(1,2) \\ 0.9U(1,1) + 0.1U(2,1) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$

[1] R. Bellman, "A markovian decision process," Technical report, DTIC Document, 1957.

Value iteration



- We can solve this problem using iterative algorithms.
- Just guess some values and keep updating U functions until it's converged.
- How to solve it iteratively?
- Is it converged?



$$U(1,1) = -0.04 + \gamma \max_a \begin{bmatrix} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1) \\ 0.9U(1,1) + 0.1U(1,2) \\ 0.9U(1,1) + 0.1U(2,1) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$

Value iteration

```
function VALUE-ITERATION(mdp,  $\epsilon$ ) returns a utility function
  inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ ,
           rewards  $R(s)$ , discount  $\gamma$ 
            $\epsilon$ , the maximum error allowed in the utility of any state
  local variables:  $U$ ,  $U'$ , vectors of utilities for states in  $S$ , initially zero
                      $\delta$ , the maximum change in the utility of any state in an iteration

  repeat
     $U \leftarrow U'$ ;  $\delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
       $U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$       (Update by Bellman EQ)
      if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
  until  $\delta < \epsilon(1 - \gamma)/\gamma$ 
  return  $U$ 
```

Value iteration


- 1st iteration

$$U(1,1) = U(1,2) = U(1,3) = U(2,1) =$$

$$U(2,3) = U(3,1) = U(3,2) = U(3,3) = U(4,1) = 0$$

$$U(4,2) = -1$$

$$U(4,3) = +1$$



$$U(1,1) = -0.04 + \gamma \max_a \begin{bmatrix} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1) \\ 0.9U(1,1) + 0.1U(1,2) \\ 0.9U(1,1) + 0.1U(2,1) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$

$$= -0.04$$

Value iteration

- 1st iteration

$$U(1,1) = U(1,2) = U(1,3) = U(2,1) =$$

$$U(2,3) = U(3,1) = U(3,2) = U(3,3) = U(4,1) = 0$$


$$U(4,2) = -1$$

$$U(4,3) = +1$$

$$U(4,3) = R(4,3) + \gamma \cdot 0 = +1$$

$$U(3,3) = -0.04 + \gamma \max_a \begin{bmatrix} 0.8U(3,3) + 0.1\underline{U(4,3)} + 0.1U(2,3) & (up) \\ 0.8U(2,3) + 0.1U(3,2) + 0.1U(3,3) & (left) \\ 0.8U(3,2) + 0.1U(2,3) + 0.1\underline{U(4,3)} & (down) \\ 0.8\underline{U(4,3)} + 0.1U(3,3) + 0.1U(3,2) & (right) \end{bmatrix} \leftarrow$$

$$= -0.04 + 0.9(0.8 * 1 + 0.1 * 0 + 0.1 * 0) = 0.68$$



Value iteration

- 2nd iteration – update U function


$$U(1,1) = U(1,2) = U(1,3) = U(2,1) =$$

$$U(2,3) = U(3,1) = U(3,2) = U(4,1) = -0.04$$

$$U(3,3) = 0.68$$

$$U(4,2) = -1$$

$$U(4,3) = +1$$



Compute $U(S)$ again until it's converged!

Value iteration

- After convergence, the robot can find optimal decision at each state

$$\pi = \arg \max_a P(s'|s, a)U(s')$$

$$= \arg \max_a \begin{bmatrix} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1) \\ 0.9U(1,1) + 0.1U(1,2) \\ 0.9U(1,1) + 0.1U(2,1) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$


$$= \arg \max_a \begin{bmatrix} 0.8 * 0.39 + 0.1 * 0.25 + 0.1 * 0.29 \\ 0.9 * 0.29 + 0.1 * 0.39 \\ 0.9 * 0.29 + 0.1 * 0.25 \\ 0.8 * 0.25 + 0.1 * 0.39 + 0.1 * 0.29 \end{bmatrix} \begin{matrix} (up) \\ (left) \\ (down) \\ (right) \end{matrix}$$

3	0.51	0.64	0.80	+1
2	0.39		0.48	-1
1	0.29	0.25	0.34	0.12
	1	2	3	4

Repeat this step for each state.


Value iteration

- After convergence, the robot can find optimal decision at each state




Value iteration


- Punishment for not arriving at goals will change the optimal policy.




$$R(s) < -1.6284$$



$$-0.4278 < R(s) < -0.0850$$



$$-0.0221 < R(s) < 0$$

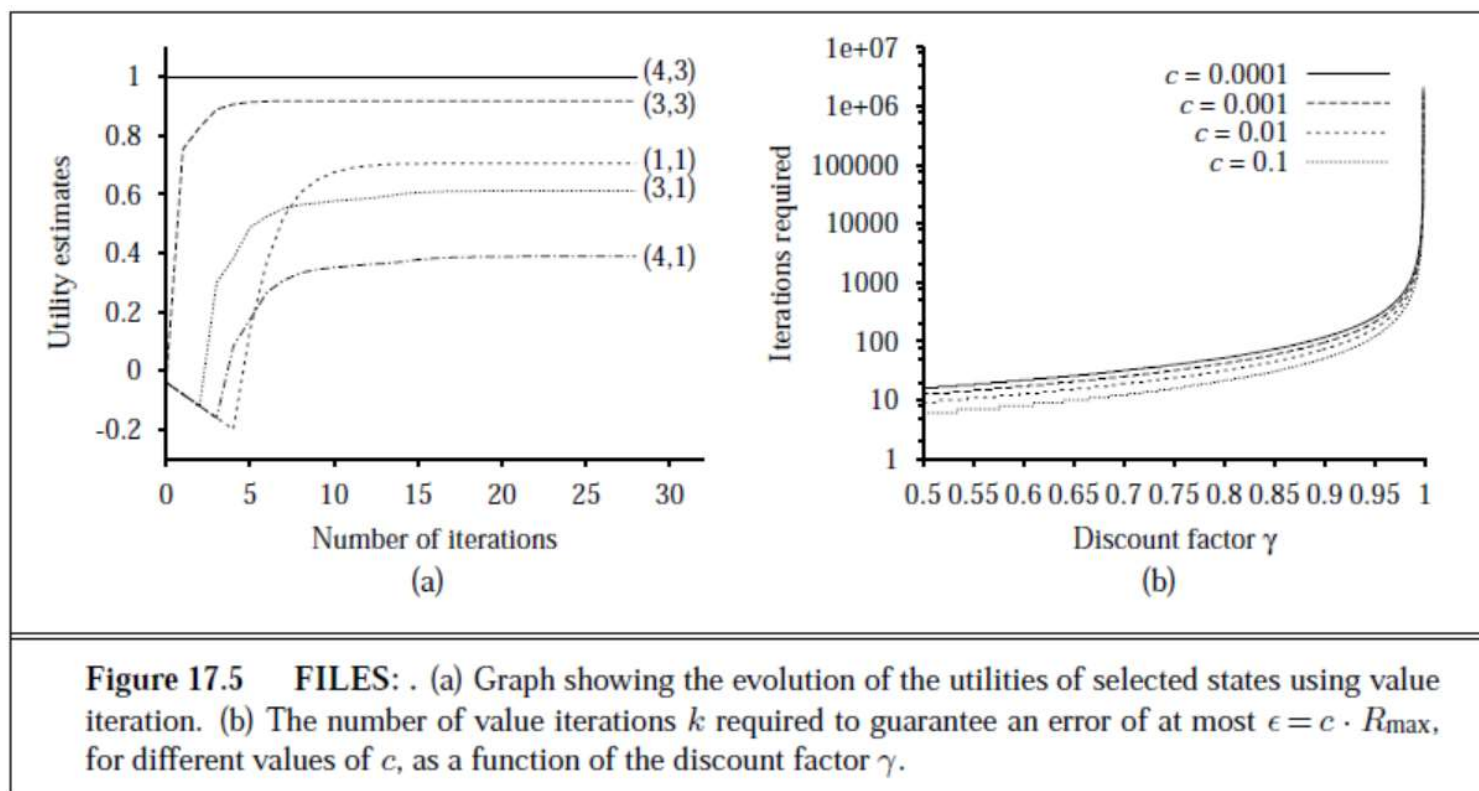


$$R(s) > 0$$

3	0.51	0.64	0.80	+1
2	0.39		0.48	-1
1	0.29	0.25	0.34	0.12
	1	2	3	4

The rewards of $S=(4,2)$ and $S=(4,3)$ will affect the optimal policy also.

Value iteration



Policy iteration

- If you check the policy of each iteration of VI algorithm, you will find the optimal policy is converged before the U function is converged.
- Hence, policy iteration could find the solution faster than value iteration.

Policy iteration

- The major concept of PI is as follows:
 - Assign actions
 - Compute U based on current actions
 - Repeat until convergence

Assign $a = \{\text{up}\}$ in each state


$$U_i(1,1) = -0.04 + 0.8U_i(1,2) + 0.1U_i(1,1) + 0.1U_i(2,1)$$

$$U_i(1,2) = -0.04 + 0.8U_i(1,3) + 0.2U_i(1,2)$$

\vdots

$$U_i(4,1) = -0.04 + 0.8U_i(4,2) + 0.1U_i(3,1) + 0.1U_i(4,1)$$

Since there is no “max” operator in PI, there are N unknown variables and N equations.
Solve $AX=B$ problem. However, it's $O(N^3)$.




Policy iteration

```
function POLICY-ITERATION(mdp) returns a policy
  inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ 
  local variables:  $U$ , a vector of utilities for states in  $S$ , initially zero
                    $\pi$ , a policy vector indexed by state, initially random

  repeat
     $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, \text{mdp})$            (Compute by Pseudo inverse)
    unchanged?  $\leftarrow$  true
    for each state  $s$  in  $S$  do
      if  $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$  then do
         $\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$    (Update policy)
        unchanged?  $\leftarrow$  false
  until unchanged?
  return  $\pi$ 
```

Policy iteration

- Results



In this initial setting, the robot got the optimal solution after 4 iterations.

Policy iteration




Figure 17.6 FILES: . The maximum error $\|U_i - U\|$ of the utility estimates and the policy loss $\|U^{\pi_i} - U\|$, as a function of the number of iterations of value iteration.


Conclusions

- MDP is a model for finding sequential optimal decisions.
 - State: fully observable
 - State transition: stochastic ([motion model](#))
- Bellman equation is the key to solve MDP problems.
- MDP is widely applied to decision problems. However, the basic assumption is that the state is deterministic!
- POMDP is a model for finding sequential optimal decisions.
 - State: stochastic ([sensor model](#) → [Bayes theorem](#))
 - State transition: stochastic ([motion model](#))

Conclusions

- LRTA*

Deterministic action



$$s, a \rightarrow s'$$

L2: Uninformed search

L3: Heuristic search (LRTA*)


L4: Adversarial search

L5: Bayes theorem

L6: Bayes theorem over time

MDP (RL)

Probabilistic actions



$$P(s'|s, a)$$

L7: MDP

L9: Reinforcement learning

L10: GP and LWPR


L11: Naïve Bayes and Perceptron

L12: Adaboost

L13: Deep learning and DRL

POMDP

Probabilistic actions and states



$$P(s'|s, a), P(s)$$

L8: POMDP

Q&A

