


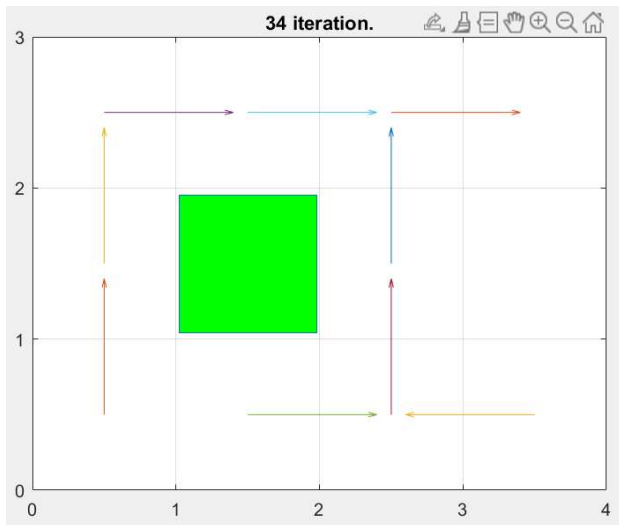
Reinforcement Learning

KUO-SHIH TSENG (曾國師)
DEPARTMENT OF MATHEMATICS
NATIONAL CENTRAL UNIVERSITY, TAIWAN

2021/04/28


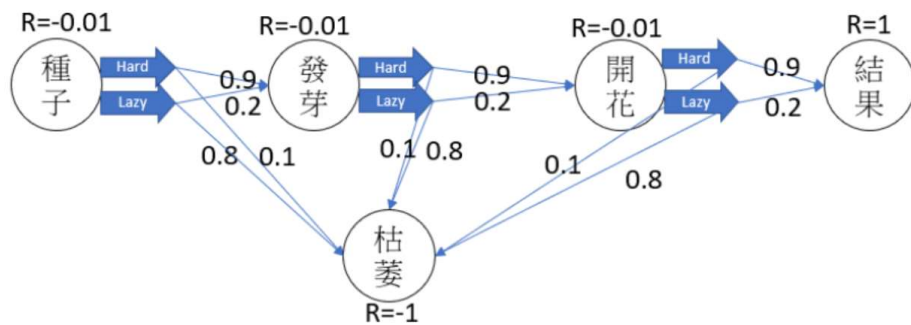
HW2

- MDP



HW2

- MDP example



Midterm

1. Perception

- a) Smoothing (20%)
- b) Bayesian search in 2 grids (15%)
- c) Bayesian search in 3 grids (15%)

2. Decision-making

- a) MDP (20%)
- b) Bellman EQ (15%)
- c) MCTS (15%)





Figure 1: Illustration of the problem.

Course Announcement

- Deadline of final project proposal: **05/12 0am**
- Final Project Proposals should include the following :
 - The goal of learning
 - MDP formulation of your project
 - State, action, transitional probability, rewards, and discount terms.
 - Features of Q function
 - The data flow of your robot
- You can find the sample files of final project on LMS.
- You should cite the prior work (Reference).

Course Announcement

- HW3 was released (Deadline: **5/26 0am**):
- A RL problem (40%)
- A supervised learning problem (40%)
- A RL proof (20%)
- You can start to work on P1 and P3 of HW3 today. Work on it **ASAP!**

Course Announcement

- The grade for this course will consist of the following components:
 - HW1 10%
 - HW2 10%
 - Midterm Exam 30%
 - HW3 10%
 - Project Proposal (1-2 page) **10%**
 - Project Presentation and Demonstration **10%**
 - Project Report (4-8 pages) **20%**
- NOTICE: I will **NOT** do curve fitting (e.g., “ $\sqrt{X} \cdot 6$ ” for your score) for your scores.

Course Announcement

- To complete your final project, we will have a ROS tutorial for Minibot on 5/5.
 - Take your laptop to class
 - Take your **Minibot** or **bebop** to class with full charged battery.
 - Install virtualbox <https://www.virtualbox.org/> before the class
- Download the image file from Minibot wiki page before the class
 - VirtualBox Image (password: hypharos) :
<https://drive.google.com/open?id=1xTVsPet6WT48Psete6ilkgg-gi1QdOht>
- You can use your Minibot or Bebop in this building. You should return your robot to M-213 when you are not using it.

Course Announcement

May 2021

Su	Mo	Tu	We	Th	Fr	Sa
	3	4	5	6	7	8
	10	11	12	13	14	15
	17	18	19	20	21	22
	24	25	26	27	28	29
	31	1	2	3	4	5
	7	8	9	10	11	12
	14	15	16	17	18	19
	21	22	23	24	25	26

ROS
Tutorial

L11 NB
Perceptron

L12
Adaboost

L10 GP

L13
DL and DRL

L14
Kmeans, EM

Final Project
Presentation

Final Project
DEMO

HW3 released

Proposal DL

HW3 DL

Final Project report


Outline

- Need
- Reinforcement learning
- Monte Carlo methods
- TD- learning
- Q-learning
 - Tubular Q-learning
 - function based Q-learning
- Linear regression
- Regularization
- Appendix – More about RL

Outline

- LRTA*

Deterministic action



$$s, a \rightarrow s'$$

L2: Uninformed search

L3: Heuristic search (LRTA*)


L4: Adversarial search

L5: Bayes theorem

L6: Bayes theorem over time

MDP (RL)

Probabilistic actions



$$P(s'|s, a)$$

L7: MDP

L9: Reinforcement learning

L10: GP and LWPR


L11: Naïve Bayes and Perceptron

L12: Adaboost

L13: Deep learning and DRL

POMDP

Probabilistic actions and states



$$P(s'|s, a), P(s)$$


L8: POMDP

Outline

[Problem solving]

Search problems


Adversarial Search



[Perception and Uncertainty]

Bayes Theorem

Bayes Filter and Smoothing






[Learning and Decision-making]

Supervised learning

Unsupervised learning

Reinforcement learning



Need

- The disadvantages of iterative approaches for solving MDP are
 - (1) The robot needs to know the environments or states.
 - (2) they need a transition probability model;
 - (3) they need to update each state to converge to a solution.
- For most online learning applications, the robots do not know P and need to update utility immediately.

[*GIVEN*]

S : *state*

A : *action*

$P(s'|s, a)$: *Transition probability*

$R(s, a)$: *reward*

γ : *discount*


[*Find*]

$$\pi^* = \arg \max U^\pi(s)$$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

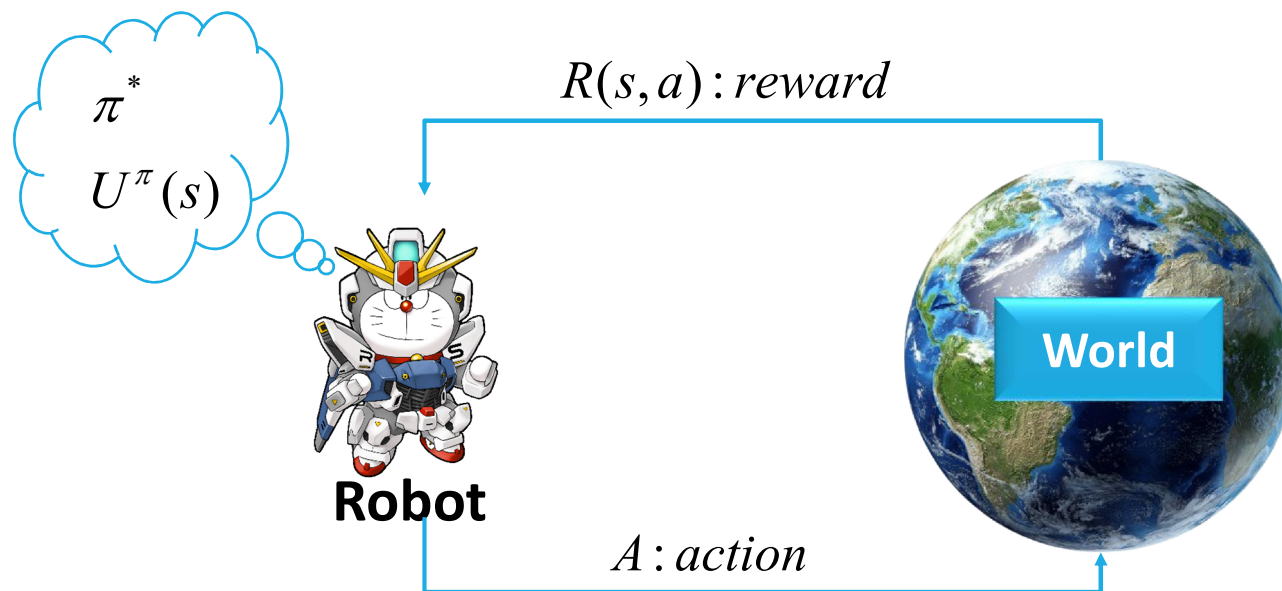
Need

- Let the robot interact with the world and receive reward immediately.
- The robot can try many times to update its utility function. Then, the robot can find optimal actions based on the learned utility function.





Reinforcement learning

- Let the robot interact with the world and receive the reward immediately. RL is inspired by animals. Animals learn the world via interactions!

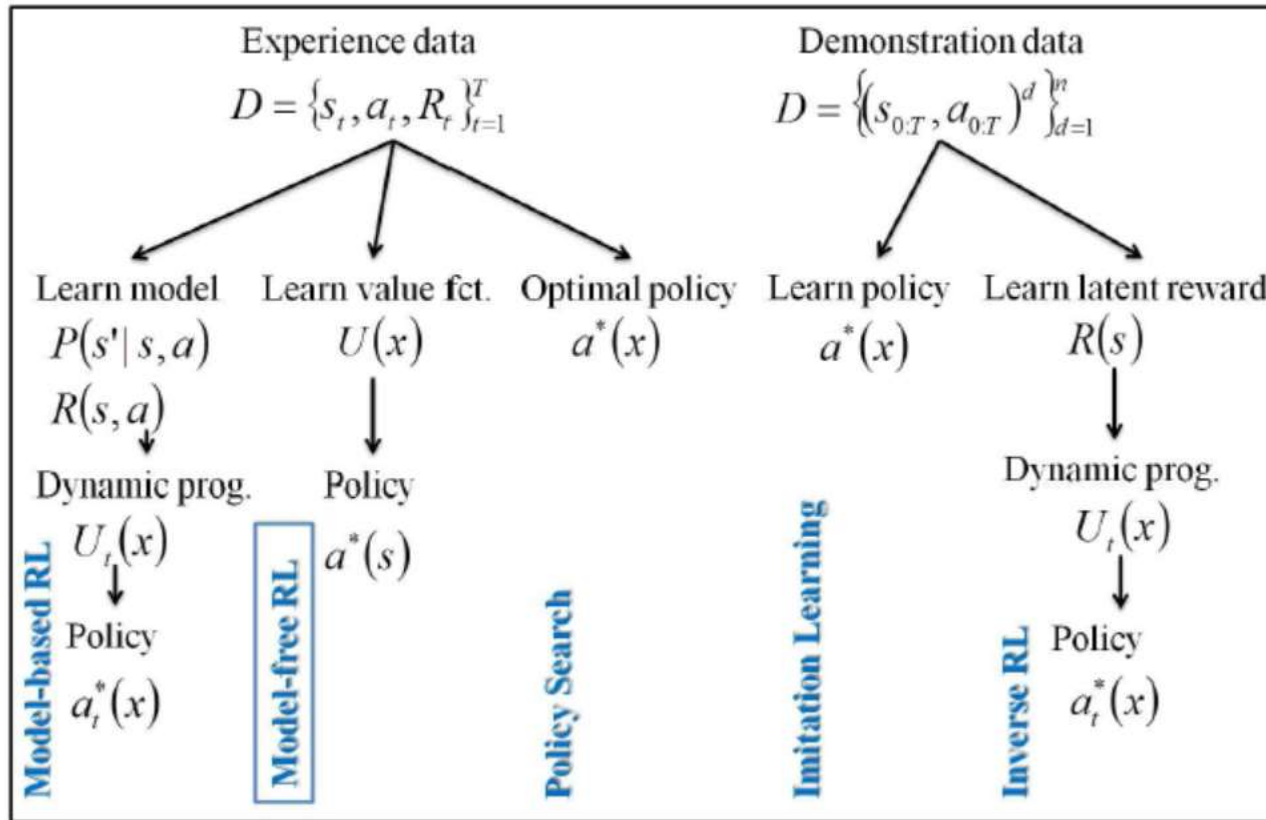


Reinforcement learning

- During the past 10 years, robots are able to do amazing **control** tasks (~60 DOFs) using RL.
- In 2015, AlphaGo can beat the best human players.
- In 2015, AI can play 49 Atari games.
- RL+DL is the future of AI?




Reinforcement learning




[1] Marc Toussaint. Machine learning and robotics. ICML tutorial, 2011.

Reinforcement learning



From: Facebook OURSKY

Reinforcement learning



Richard S. Sutton and Andrew G. Barto, "Reinforcement Learning: An Introduction," MIT Press, 2nd edition, 2018.

Monte Carlo methods

- Monte Carlo approaches can be applied to compute the utility function.
- For example, there are some trial samples.



$(1, 1) \rightsquigarrow (1, 2) \rightsquigarrow (1, 3) \rightsquigarrow (1, 2) \rightsquigarrow (1, 3) \rightsquigarrow (2, 3) \rightsquigarrow (3, 3) \rightsquigarrow (4, 3) + 1$
 $(1, 1) \rightsquigarrow (1, 2) \rightsquigarrow (1, 3) \rightsquigarrow (2, 3) \rightsquigarrow (3, 3) \rightsquigarrow (3, 2) \rightsquigarrow (3, 3) \rightsquigarrow (4, 3) + 1$
 $(1, 1) \rightsquigarrow (2, 1) \rightsquigarrow (3, 1) \rightsquigarrow (3, 2) \rightsquigarrow (4, 2) - 1$

- The utility function can be solved by Bellman EQ and these samples:

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

or

$$U(s) = R(s) + \gamma \max_a \left[\sum_{s'} U(s') P(s' | s, a) \right]$$



TD-learning

- Temporal-difference (TD) learning is proposed based on dynamic programming (DP) and Monte Carlo concepts.
- TD updates the utility function based on the difference of utility after each action. Hence, it needs a lot of samples (Monte Carlo) and updates parts of utility (DP) without reaching terminals.
- The updated equation of the utility function is as follows:


$$U(s) \leftarrow U(s) + \alpha(R(s) + \gamma U(s') - U(s))$$

$U(s)$: utility function

$R(s)$: reward

α : learning rate

- TD is a **model free** approach.



TD-learning

- Given an action (a), reward (R) and next state (s'), the robot updates $U(s)$.

$$s = (3,1)$$

$$s' = (3,2)$$

$$\gamma = 1, \alpha = 0.1$$

$$U(s) \leftarrow U(s) + \alpha(R(s) + \gamma U(s') - U(s))$$

$$U(s) \leftarrow 0.34 + 0.1(-0.04 + 0.48 - 0.34)$$

$$U(s) = 0.35$$

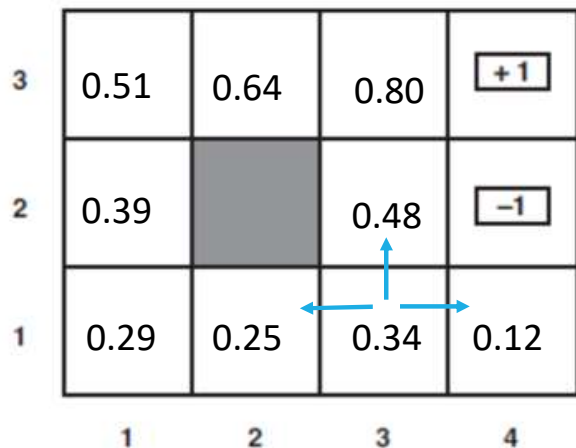
When $(R(s) + \gamma U(s') - U(s)) = 0$,

$U(s)$ will not be updated.

3	0.51	0.64	0.80	+1
2	0.39		0.48	-1
1	0.29	0.25	0.34	0.12
	1	2	3	4

TD-learning

- A disadvantage of TD is that it needs a lot of samples since the robot does not know **what the next action is** while collecting data.
- To improve TD, the robot needs a function with **actions** and **states** variables. $U(s) \rightarrow Q(s,a)$



Q-learning

- In Q-learning, the robot learns an **action-utility** function $Q(s,a)$.
- With the learned Q function, the robot knows what the next best action is. Since Q-learning does not need transition probability and action section models, it is one of the *model-free* methods for reinforcement learning.


3	0.51	0.64	0.80	+1
2	0.39		0.48	-1
1	0.29	0.25	0.34	0.12
	1	2	3	4

$s = (3,1)$

$$\begin{cases} Q(s, a = \text{left}) \\ Q(s, a = \text{up}) \\ Q(s, a = \text{right}) \end{cases}$$

Pick up the action with the maximal Q value

$$U(s) = \max_a Q(s, a)$$



Q-learning

- Bellman equation: Assuming the agent chooses the optimal action, The utility of a state is the immediate reward for that state + the expected discounted utility of the next state.

$$U(s) = R(s) + \gamma \max_a \left[\sum_{s'} U(s') P(s' | s, a) \right]$$


immediate reward expected discounted utility

- Bellman equation of Q-learning

$$Q(s, a) = R(s) + \gamma \left[\sum_{s'} \max_{a'} Q(s', a') P(s' | s, a) \right]$$

immediate reward expected discounted utility

$$Q(s, a) = R(s) + \gamma \max_{a'} Q(s', a')$$



Q-learning

- The TD learning can be extended to Q learning.

$$U(s) \leftarrow U(s) + \alpha(R(s) + \gamma U(s') - U(s))$$

TD learning

$$\because U(s) = \max_a Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

Q learning

or

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left(R(s) + \gamma \max_{a'} Q(s', a') \right)$$

$Q(s, a)$: action – utility function (Q function)

$R(s)$: reward

α : learning rate

Q-learning

- Q-learning

[*GIVEN*]

S : *state*

A : *action*

$R(s, a)$: *reward*

γ : *discount*

α : *learning rate*

[*Find*]

$Q(s, a)$: action – utility function (Q function)

$$U(s) = \max_a Q(s, a)$$


$$\pi^* = \arg \max U^\pi(s)$$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

Q-learning

- Trade-off between exploration and exploitation of the Q function.
- If the robot always select the action with the maximal utility, it could follow the same path without exploring other paths.
- ϵ -greedy was proposed to solve this issue. The robot will action randomly with ϵ probability and select best action with $(1 - \epsilon)$ probability.



Exploration:
Random actions


or

$$\begin{array}{c}
 Q(s, a = up) \\
 \uparrow \\
 Q(s, a = left) \quad \quad Q(s, a = right)
 \end{array}$$


Exploitation:
Pick up the best action
based on Q

Tubular Q-learning

- Although $Q(s,a)$ can provide the utility of each action at each state, it means there are more values to learn!



$U(s)$: There are 9 values to learn.

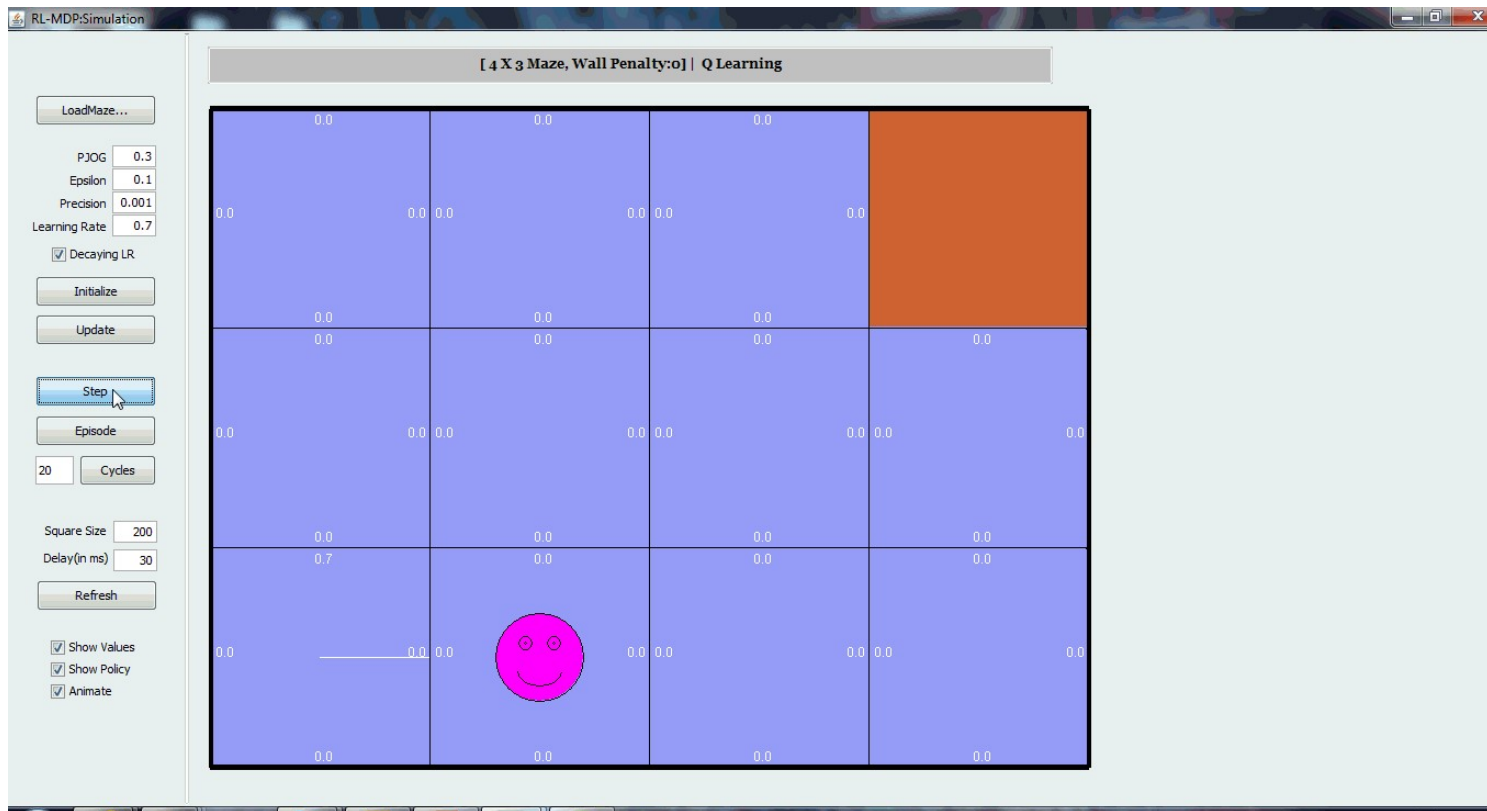


$Q(s,a)$: There are $9 \cdot 3$ values to learn.

There are $|S| \cdot |A|$ values in Q table

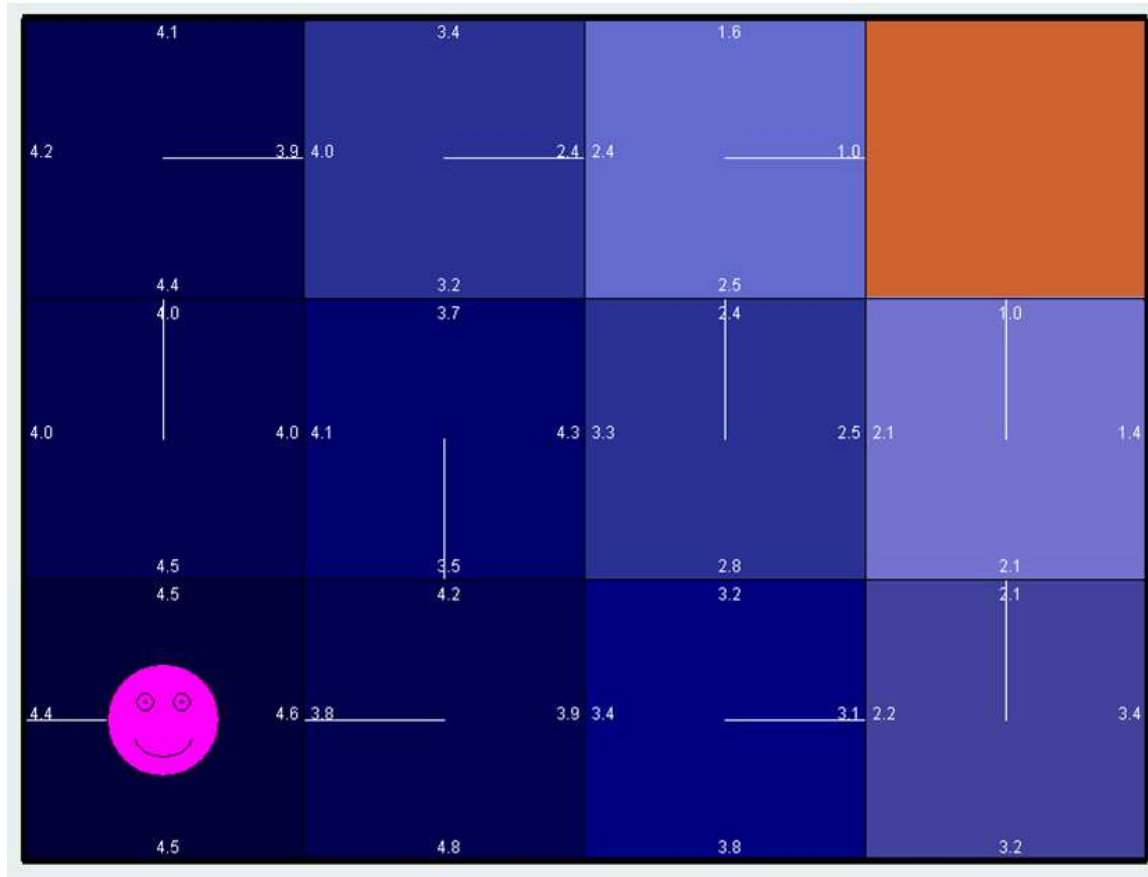
$|S|$: # of states, $|A|$: # of actions

Tubular Q-learning



<http://www.cs.cmu.edu/~awm/rlsim/>

Tubular Q-learning



This is similar to HW3 P1.

Tubular Q-learning

- Disadvantage of tubular Q-learning
 1. Storage space: $|S|X|A|$
 2. Exploration time
 3. Convergence time
- If the state space and action space is larger, the Q function is difficult to converge.
- To solve these issues, function-based Q learning was proposed. The Q function can be approximated by functions.

Function-based Q learning

- To improve tubular Q learning, function approximation was adopted.
- There are two kinds of approximations
- Parametric approximation:
 - The parametric approximators are mappings from a parameter space to the space Q functions. For example, polynomial basis (feature) is a popular approach.
- Nonparametric approximation:
 - The non parametric approximators are derived from the available data. For example, Gaussian processes (GP) and neural network (NN) are popular after 2005.

Function-based Q learning

- Parametric approximation:
- The Q function is approximated by some features.

$$Q(s, a) = \sum_i w_i f_i(s, a)$$

w : weighting vector

f : features

- How to find w ? \rightarrow gradient descent
- How to find f ? \rightarrow hand crafting

Function-based Q learning

- There are two ways to compute W :
- Offline approach: least square (pseudo inverse)

$$\min [Q(s, a) - \mathcal{Q}(s, a)]^2 = \min [(Y - WF)^T (Y - WF)]$$

$$W = (F^T F)^{-1} F^T Y$$

- Online approach: gradient descent (online least-squares)

$$E = \frac{1}{2} [(Y - WF)^T (Y - WF)] \Rightarrow w_i \leftarrow w_i - \alpha \frac{\partial E}{\partial w_i}$$

$$w_i \leftarrow w_i + \alpha \left[R(s) + \gamma \max_{a'} \mathcal{Q}(s', a') - \mathcal{Q}(s, a) \right] \frac{\partial \mathcal{Q}(s, a)}{\partial w_i}$$

Function-based Q learning

- Online approach: gradient descent (online least-squares)

$$w_i \leftarrow w_i + \alpha \left[R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \frac{\partial Q(s, a)}{\partial w_i}$$


EX :

$$Q(s, a) = w_0 + w_1 x + w_2 y$$

$$w_0 \leftarrow w_0 + \alpha \left[R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

$$w_1 \leftarrow w_1 + \alpha \left[R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] x$$

$$w_2 \leftarrow w_2 + \alpha \left[R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] y$$





Function-based Q learning

- How to find good features?
 - The distance between the robot and the goal (4,3)
 - Polynomial terms of x and y


$$Q(s, a) = \sum_i w_i f_i(s, a)$$

$$= \underline{w_0} + \underline{w_1}x + \underline{w_2}y + w_3 \sqrt{(x - x_g)^2 + (y - y_g)^2}$$



Function-based Q learning

- How to find good features for Pac-Man?
 - Distance to closest ghost
 - Distance to closest dot
 - Number of ghosts
 - $1 / (\text{dist to dot})^2$
 - ...etc.



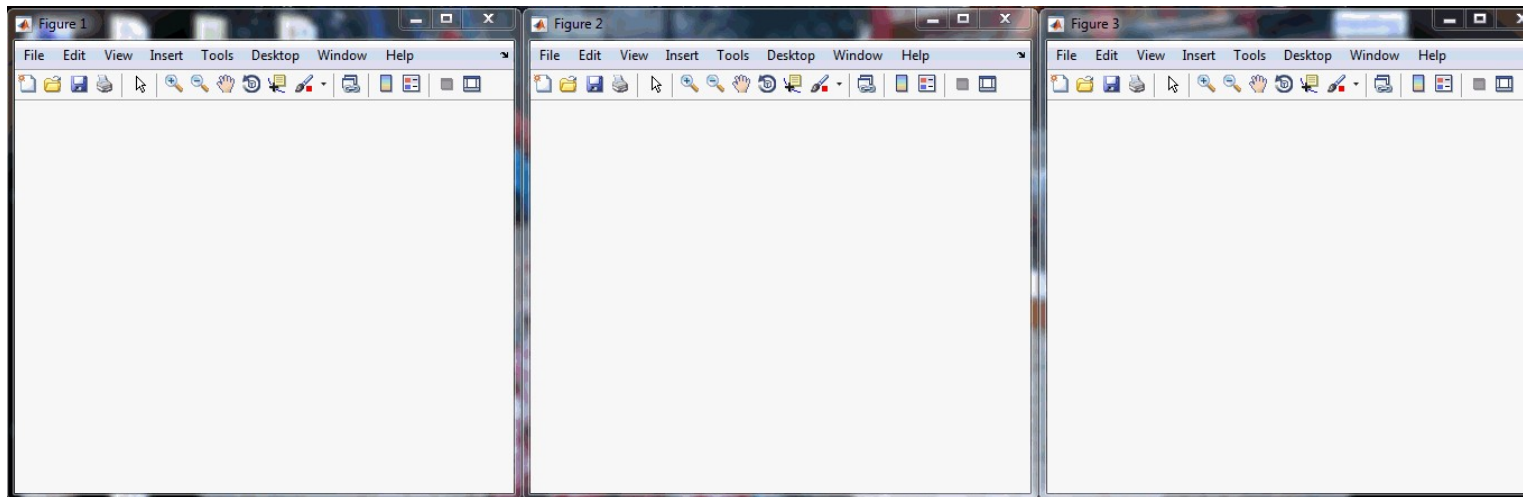
$$Q(s, a) = w_0 + w_1 \sqrt{(x - x_g)^2 + (y - y_g)^2} \\ + w_2 \sqrt{(x - x_d)^2 + (y - y_d)^2} \\ + w_3 \left(1 / ((x - x_d)^2 + (y - y_d)^2) \right)$$

Function-based Q learning

- Q-learning using 5 features

$$Q(s, a) = \sum w_i f_i$$

- Basis functions
 - Constant
 - Distance
 - Delta_theta
 - UPPER
 - Turn



Function-based Q learning

- Function-based Q learning works but humans need to find features (hand crafting).
- Finding good features is the key to function-based Q learning.
- As aforementioned examples, we need domain knowledge to find good features. This processing is called feature engineering.
- The features worked well for A problem could not work for B problem.
- Hence, researchers also adopted nonparametric approximation (GP and NN) for Q learning.

Function-based Q learning

- Q learning (Tabular v.s. function approximation)

1. run $a = \max_a Q(s, a)$

2. get reward $R(s, a)$

3. $Q(s, a) \leftarrow Q(s, a) + \alpha \left[R(s, a) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$

- ☹ Storage space
- ☹ Exploration time
- ☹ Convergence time


$$Q(s, a) = \sum w_i f_i$$

1. run $a = \max_a Q(s, a)$

2. get reward $R(s, a)$

3. $w_i \leftarrow w_i + \alpha \left[R(s, a) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] f_i$

- 😊 Storage space ↓
- 😊 Exploration time ↓
- 😊 Convergence time ↓
- 😊 ML techniques
- ☹ Find features?



Function-based Q learning

- Function-based Q learning works but humans need to find features (hand crafting).
- Issues of RL:
 - Credit assignment issues
 - Exploration and exploitation

Linear regression

- The principles of the theory are derived, as are those of rational mechanics, from a very small number of primary facts.

– Joseph Fourier, 1878

$$y = w_0 \bullet 1 + w_1 \bullet \cos(\omega_1 x) + w_2 \bullet \sin(\omega_2 x)$$

$$y = w_0 \bullet 1 + w_1 \bullet x + w_2 \bullet x^2 + \dots$$

$$y = w_0 \bullet 1 + w_1 \bullet f_1(x) + w_2 \bullet f_2(x) + \dots$$

- These functions are called “basis.” Linear regression is to find the weighting vector.


Linear regression

- A popular approach is to minimize the least square: E





$$E = (y - \hat{y})^T (y - \hat{y}) = (y - WF)^T (y - WF)$$

$$\min [Q(s, a) - Q(s, a)] = \min [(Y - WF)^T (Y - WF)]$$

$$W = (F^T F)^{-1} F^T Y$$



- Offline approach: least square (pseudo inverse)
- Online approach: gradient descent (online least-squares)
- Another issue is overfitting. The regression try to fit the training data.



Regularization

- Regularization is a way to avoid overfitting issues.

$$\text{cost}(h) = \text{Loss}(h) + \lambda \text{Complexity}(h)$$

$$\text{Complexity}(h_w) = L_q(\mathbf{w}) = \sum_i |w_i|^q$$

- For example,



$$\min \left[(Y - WX)^2 + \lambda |W|^2 \right] \dots \text{L2 norm}$$

$$\min \left[(Y - WX)^2 + \lambda |W| \right] \dots \text{L1 norm}$$

- These two minimization problems can be solved by gradient descent approaches.

Regularization


- L1 norm and L2 norm.
- L1 norm has the sparse property! L1 norm is applied to ***sparse learning***, which is to learn the weighting and **enforce** elements of weighting vector are 0. (Feature selection!)



Regularization

- Coverage function approximation
- Given: (X_i, Y_i) , $i=1 \sim m$, $X_i = \{x_i, y_i\}$
- Find: W

robot1 robot2



X : Robot position, Y : Coverage

$$Y_{(m,1)} = X_{(m,n)} W_{(n,1)}$$

$$Y = w_0 1 + w_1 x + w_2 x^2 + w_3 x^3 + w_4 xy + w_5 xy^2 + w_6 x^2 y + w_7 y + w_8 y^2 + w_9 y^3$$




$$W = (X^T X)^{-1} X^T Y \quad \text{1. Pseudo inverse} \rightarrow \text{overfitting?}$$

or

$$\mathbf{w}_{k+1} = \tau_{\lambda t} \left(\mathbf{w}_k - 2t \mathbf{X}^T (\mathbf{X} \mathbf{w}_k - \mathbf{Y}) \right) \quad \text{2. ISTA (L1 norm)} \rightarrow \text{feature selection}$$

Iterative Shrinkage-Thresholding algorithm

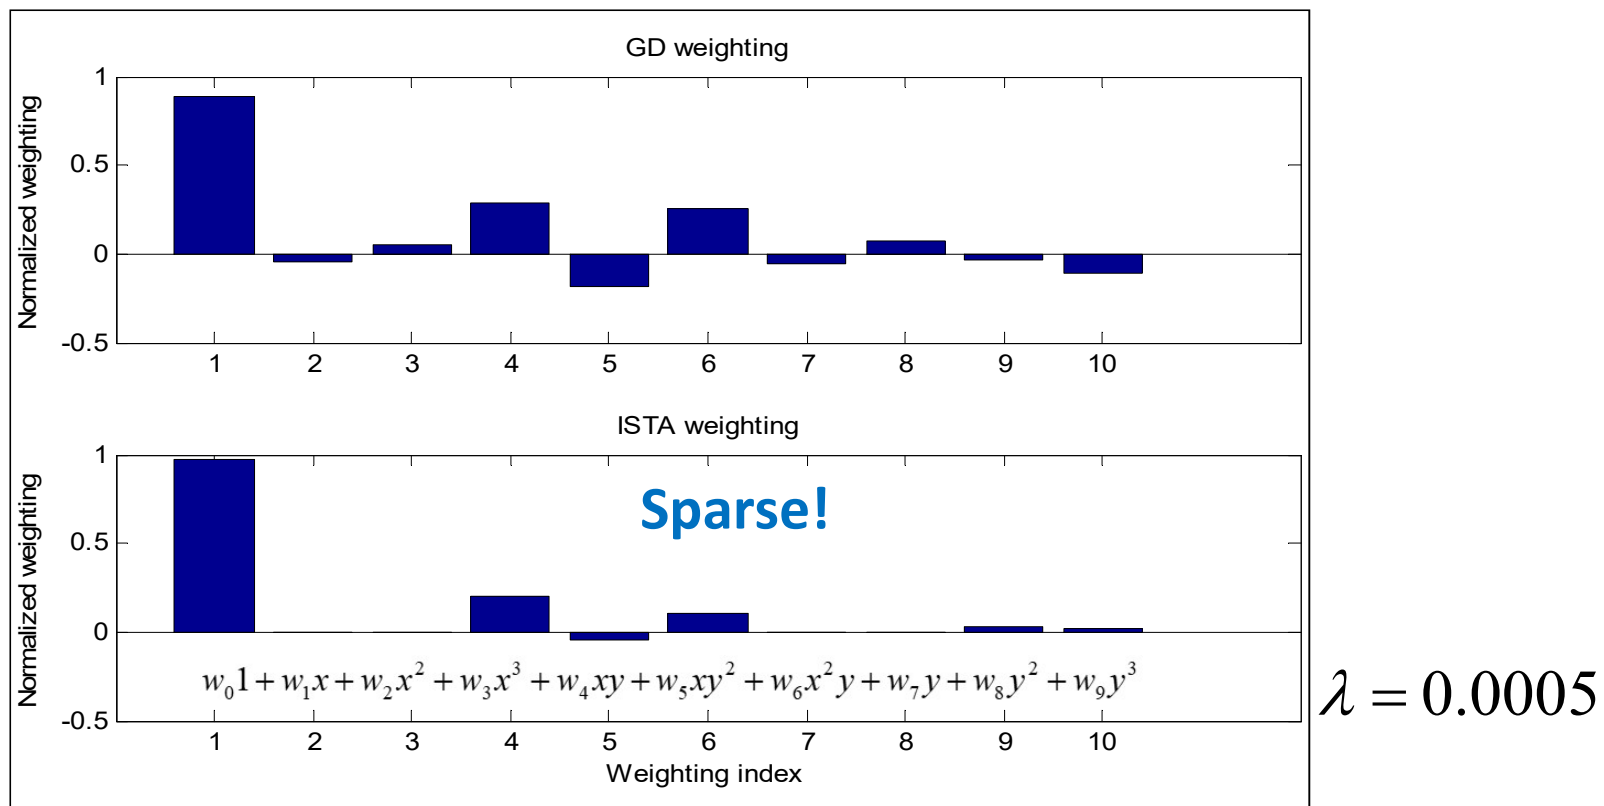
Regularization



Overlapping line

Regularization


- Results:




Regularization

- Results:

Approximation using Gradient Descent



Approximation using ISTA



Q&A



Appendix – More about RL

- On-policy (sarsa) v.s. off-policy
- Policy gradient
- Actor-critic

Appendix

- “You don’t learn to walk by following rules. You learn by doing, and by falling over.”
— Richard Branson, Entrepreneur